Bachelor Thesis

# Benchmark of RISC-V in BTOR2

## Jan Krister Möller

### Examiner: Dr. Mathias Fleury

University of Freiburg

Faculty of Engineering

Department of Computer Science

Chair of Computer Architecture

September 24, 2025

# Declaration

I hereby declare that I am the sole author and composer of my thesis and that no other sources or learning aids, other than those listed, have been used. Furthermore, I declare that I have acknowledged the work of others by providing detailed references of said work.

I hereby also declare that my Thesis has not been prepared for another examination or assignment, either wholly or excerpts thereof.

_____          _____

Place, Date                       Signature

# Declaration on Usage of Generative AI

I hereby declare that, with the approval of my examiner, I have employed the generative AI tool "GitHub Copilot" during the preparation of this thesis solely for spellchecking and enhancing the formality of my written expression. Furthermore, I expressly confirm that this tool was not used to generate data or as a source of factual information or content for this thesis.

_____          _____

Place, Date                                                    Signature

# Abstract

(TODO: write abstract)

# Contents

# List of Figures

# List of Tables

# List of Algorithms

# 1 Introduction

RISC-V, an open-source instruction set architecture, has gained significant attention due to its flexibility and extensibility and was even mentioned by the European Union for broader adoption [1]. Ensuring the correctness of RISC-V implementations is therefore of great importance, particularly in safety-critical applications.

Model checking has emerged as a powerful method for formally verifying hardware and software systems. Among the available model checking formats, BTOR2 has become a widely adopted standard for word-level hardware verification. As RISC-V is a processor architecture, it is a fitting choice. This thesis investigates the feasibility and effectiveness of applying model checking to RISC-V using the BTOR2 format.

The primary objective of this work is to develop tools that translate RISC-V processor states into BTOR2 models, enabling systematic verification, and to evaluate the performance of the model. This is extended by a set of tools to check if the model functions correctly. By implementing a suite of test cases and analyzing the results across different model checkers, this thesis aims to provide insights into the suitability of BTOR2-based model checking for RISC-V and to identify potential limitations and advantages of this approach.

The following chapters present the theoretical foundations, the methodology for transforming RISC-V states to BTOR2, the benchmarking process, and a discussion of the results.

# 2 RISC-V

As the first foundation for my benchmarks and, consequently, this thesis, I will discuss RISC-V and its operational principles.

## 2.1 Overview

RISC-V is an open-source instruction set architecture (ISA) first published in May 2011 by A. Waterman et al. [2]. As indicated by its name, it is based on the RISC design philosophy. RISC stands for Reduced Instruction Set Computer and its concept is to only include a small set of elemental and easy to execute instructions. With this, the decoding and execution of instructions can be faster compared to CISC (Complex Instruction Set Computer) design philosophy.

Since 2015, the development of RISC-V has been coordinated by the RISC-V International Association, a non-profit corporation based in Switzerland since 2020 [3]. Its objectives include providing an *open* ISA that is freely available to all, a *real* ISA suitable for native hardware implementation, and an ISA divided into a *small* base integer ISA usable independently, for example in educational contexts, with optional standard extensions to support general-purpose software development [4, Chapter 1].

Currently, RISC-V comprises four base ISAs: RV32I, RV64I, RV32E, and RV64E, which can be extended with one or more of the 47 ratified extension ISAs [4, Preface].

For the purposes of this work, I focus on a subset of the RV64I ISA.

## 2.2 The RV64I ISA

RV64I is not overly complex, but its structure is essential for understanding the subsequent work presented in this thesis. Therefore, I will explain all elements relevant to my work.

RV64I features 32 64-bit registers, labeled $x0$–$x31$, where $x0$ is hardwired to zero across all bits. Registers $x1$–$x31$ are general-purpose and may be interpreted by various instructions as collections of booleans, two's complement signed binary integers, or unsigned integers. Additionally, there is a non-accessible register called $pc$, which serves as the program counter and holds the address of the current instruction [4, Chapters 4.1, 2.1].

In RV64I, memory addresses are 64 bits in size. As the memory model is defined to be single-byte addressable, the address space of RV64I encompasses $2^{64}$ bytes [4, Chapter 1.4]. The format of the memory is little endian, so the lower bits of a number are placed at lower addresses.

Like nearly all standard ISAs of RISC-V, RV64I employs a standard instruction encoding length of 32 bits, or one *word*. Only the compressed extension named C introduces instructions with a length of 16 bits [4, Chapter 1.5], but this special case is not considered here. All RV64I instructions are encoded in one of the six formats illustrated in Figure 1. These formats may consist of

- The *opcode*:
  The opcode is used to differentiate between groups of instructions. It also defines the format type of the instruction.

- *rd*:
  This is the destination register.

**Figure 1:** RV64I encoding formats, used in [4, Chapter 2.3]

- *funct*3:

  This is used to differentiate between instructions with the same *opcode*.

- *rs*1 & *rs*2:

  These are the source registers.

- *funct*7:

  This is used for further distinctions between instructions if there are more than eight instructions in an opcode group and *funct*3 does not suffice.

- *imm*:

  This is an immediate value. In square brackets after *imm* is designated a subfield of the immediate which is represented by these bits. From these subfields, non-defined lower bits are filled with zeros whereas the highest defined bit is sign-extended to fill all non-defined higher bits.

The design of these formats results in the following features:

- Due to RISC-V's little-endian nature, the *opcode*, which encodes the general instruction, is always read first. Further specification of the instruction via *funct*3 and *funct*7 is consistently located at the same positions.

- If utilized by the instruction, $rd$, $rs1$, and $rs2$ are also always found in the same locations, simplifying decoding.

- The highest bit of $imm$ is always bit 31, making it straightforward to sign-extend the immediate value.

The instructions relevant to my work are listed in Table 1 I have divided the instructions in Table 1 into nine groups based on their operations.

LUI and AUIPC move a high immediate into $rd$. In the case of AUIPC, the $pc$ is added to this value. JAL and JALR instructions are unconditional jumps, where for JAL $imm$ is added to $pc$ and for JALR $imm$ is added to $rs1$ and set as $pc$. Both link to the next instruction (current $pc + 4$) in $rd$.

*branch* instructions are conditional jumps. $rs1$ is compared to $rs2$ and if the comparison holds, $imm$ is added to $pc$. The comparisons are $=$ for BEQ, $\neq$ for BNE, $<$ for BLT, and $\geq$ for BGE. In these instructions, the values in $rs1$ and $rs2$ are handled as two's complement integers. The suffix *U in an instruction generally designates an unsigned operation. In this case, the values in $rs1$ and $rs2$ are handled as unsigned integers. Apart from this, they work as their counterpart without the suffix.

*load* instructions load values from memory at address $(rs1 + imm)$ into $rd$, either at Byte, Halfword, Word, or Doubleword length. By default, the value is sign-extended, and the suffix *U designates the loading of a non-sign-extended value. Conversely, *store* instructions write values from $rs2$ at the address $(rs1 + imm)$ to memory. Here also the distinction between the different lengths is made, and the lowest byte, halfword, word, or the whole doubleword is stored at the address.

All further instructions can be seen as generic operations, differentiated by their suffixes. To simplify the explanation process, all operations without any suffix and their behavior are listed in Table 2. This is almost exactly the group with opcode *op*, except the SLTU instruction, which is not suffix-free. However, as with all other

| Instr | opcode | Type | Instr | opcode | Type | Instr | opcode | Type |
|---|---|---|---|---|---|---|---|---|
| LUI | *lui* | U | SB | *store* | S | ADD | *op* | R |
| AUIPC | *auipc* | U | SH | | | SUB | | |
| JAL | *jal* | J | SW | | | SLT | | |
| JALR | *jalr* | I | SD | | | SLTU | | |
| BEQ | *branch* | B | ADDI | *op-imm* | I | XOR | | |
| BNE | | | SLTI | | | OR | | |
| BLT | | | SLTIU | | | AND | | |
| BGE | | | XORI | | | SLL | | |
| BLTU | | | ORI | | | SRL | | |
| BGEU | | | ANDI | | | SRA | | |
| LB | *load* | I | SLLI | | I* | ADDW | *op-32* | R |
| LH | | | SRLI | | | SUBW | | |
| LW | | | SRAI | | | SLLW | | |
| LD | | | ADDIW | *op-imm-32* | I | SRLW | | |
| LBU | | | SLLIW | | I** | SRAW | | |
| LHU | | | SRLIW | | | | | |
| LWU | | | SRAIW | | | | | |

**Table 1:** Subset of RV64I instructions

instructions with the unsigned suffix, it behaves as its signed counterpart except for handling both $rs1$ and $rs2$ as unsigned integers.

These operations can be extended by the *I suffix, which is designated by the opcode $op-imm$. This replaces $rs2$ with $imm$ in the behavior. Again, SLTI can be extended to an unsigned version SLTIU, which behaves as expected. A SUBI instruction does not exist as it is redundant; its behavior can be achieved by using ADDI with a negative immediate.

Additionally, the operations ADD, SUB, SLL, SRL, and SRA can be extended with the *W suffix. This forms the group with the opcode $op-32$. In contrast to the base instructions, these new ones behave as if the registers are only 32 bits. The result is placed in the low 32 bits of $rd$ and sign-extended to the full 64 bits. Overflows are ignored.

The last group is the combination of both suffixes *IW with the opcode $op-imm-32$. The behavior differs from the base instructions, as expected, by a replacement of $rs2$

| Instr | Behavior |
|------:|----------|
| ADD | $rd := rs1 + rs2$ |
| SUB | $rd := rs1 - rs2$ |
| SLT | $rd := 1$ if $rs1 < rs2$ else $rd := 0$ |
| XOR | $rd := rs1 \oplus rs2$, bitwise |
| OR | $rd := rs1 \vee rs2$, bitwise |
| AND | $rd := rs1 \wedge rs2$, bitwise |
| SLL | $rd := rs1$ shifted left by $rs2$, new bits are zeros |
| SRL | $rd := rs1$ shifted right by $rs2$, new bits are zeros |
| SRA | $rd := rs1$ shifted right by $rs2$, sign extend |

**Table 2:** All suffix free operations in RV64I and their behavior. All values are handled either bitwise or as signed twos complement integers

with $imm$ and only operating on 32 bits. Again, a SUBIW instruction is redundant as a negative immediate with ADDIW achieves the same result.

Compared to the full RV64I ISA, I have omitted the FENCE, ECALL, and EBREAK instructions, as without I/O interaction or an environment such as an OS or a debugger, these are not required.

For each of the instructions in Table 2, I also included the format in which each instruction is encoded. Most should be not surprising as they fit the description of the instructions. Only SLLI, SRLI, SRAI with I* and SLLIW, SRLIW, SRAIW with I** should need clarification. Both are essentialy the I format but with extra constraints. For I* the highest realistic shift amount for 64 bit registers is also 64. So the bits [11,9:6] of $imm$ have to be 0. The bit [10] gets a special role as it is used to differentiate between the two types of right shift. With I**, with the word suffix, the maximum shift amount is only 32, so the bit [5] of $imm$ must also be 0.

## 2.3 Simulation of RISC-V

To run RISC-V code, the obvious way would be to run it on a RISC-V processor. As this is not practical in my case because I do not have one, I have to simulate the

execution of RISC-V. For this, I will use a RISC-V-simulator of this described subset of RV64I in C that I have written for my bachelors project [5]. It will be used to test the BTOR2 model I will present in Chapter 4. Alternatively qemu [6] or QtRVSim [7] could be used, but QtRVSim has more detail than needed and qemu can not be fed with an inital state but only a full program. So I used my own, which takes the state of a RISC-V processor as an input as described in Section 2.3.3.

First, I explain the structure to represent a simple RISC-V processor I implemented.

## 2.3.1 Representing the State of a RISC-V Processor

The state requires a representation for all registers. $pc$ is defined as a 64-bit integer, and the other 32 registers are implemented as an array, allowing each register to be referenced by its number. Additionally, I implemented an array of flags, one for each register, to differentiate between initialized and non-initialized registers. The memory is built from single memory cells, each holding an address and its byte of content. These are accumulated in a hash table called "memorytable", hashing on the address. If adding a new cell causes a collision, it is appended to the cells already in the bucket, forming a linked list. These structures are shown in Figure 2.

## 2.3.2 Running an Instruction

After fetching the current instruction from the hash table, it must be decoded. The easiest way to decode the operation corresponding to the current instruction is a decision tree like in Figure 3. First, I mask out the opcode and match it over all implemented opcodes. From there, either this is an endpoint and the instruction is identified, or $funct3$ must be masked and matched. A final differentiation over $funct7$ might be needed, but after this every leaf in the tree coincides with an instruction. Also, with knowing the opcode of the current instruction, I know the instruction format (Figure 1). This means that I can now also extract relevant register numbers

```
typedef struct memory_cell
{                                       typedef struct memory_table
    uint64_t address;                   {
    uint8_t content;                        memory_cell *memory[TABLESIZE];
    struct memory_cell *next_cell;          uint64_t initialised_cells;
} memory_cell;                          } memory_table;

                    typedef struct state
                    {
                        uint64_t pc;
                        uint64_t regs_values[32];
                        bool regs_init[32];
                        memory_table *memory;
                    } state;
```

**Figure 2:** State representation of a RISC-V processor in the simulation [5]

and, if it exists, the immediate. The best way to get these values is to apply a mask and shift this result to its correct place. For the immediate, as it possibly is divided into multiple fields, it might be needed to add multiple partial immediates together. At this point all information in the instruction is decoded and the current state can be modified according to the operation corresponding to the leaf reached after going through the tree.

### 2.3.3 Saving the State of a RISC-V Processor

To preserve the current state of a RISC-V processor, both the registers and memory must be stored. For this purpose, I have devised the format shown in Figure 4. The RISC-V simulation uses this format as input and output. The minimal file consists only of the two designators "REGISTERS:" and "MEMORY:" and one empty line between them. Under "REGISTERS:", all registers can be listed with their corresponding value. Of course, x0 cannot be different from 0. I included the option to reference it nonetheless to have the complete state included. Under "MEMORY:", after giving an address, the memory can be filled with 1-, 2-, 4-, or 8-byte sized memory content. The given address is the starting address of the content. As RISC-V

10
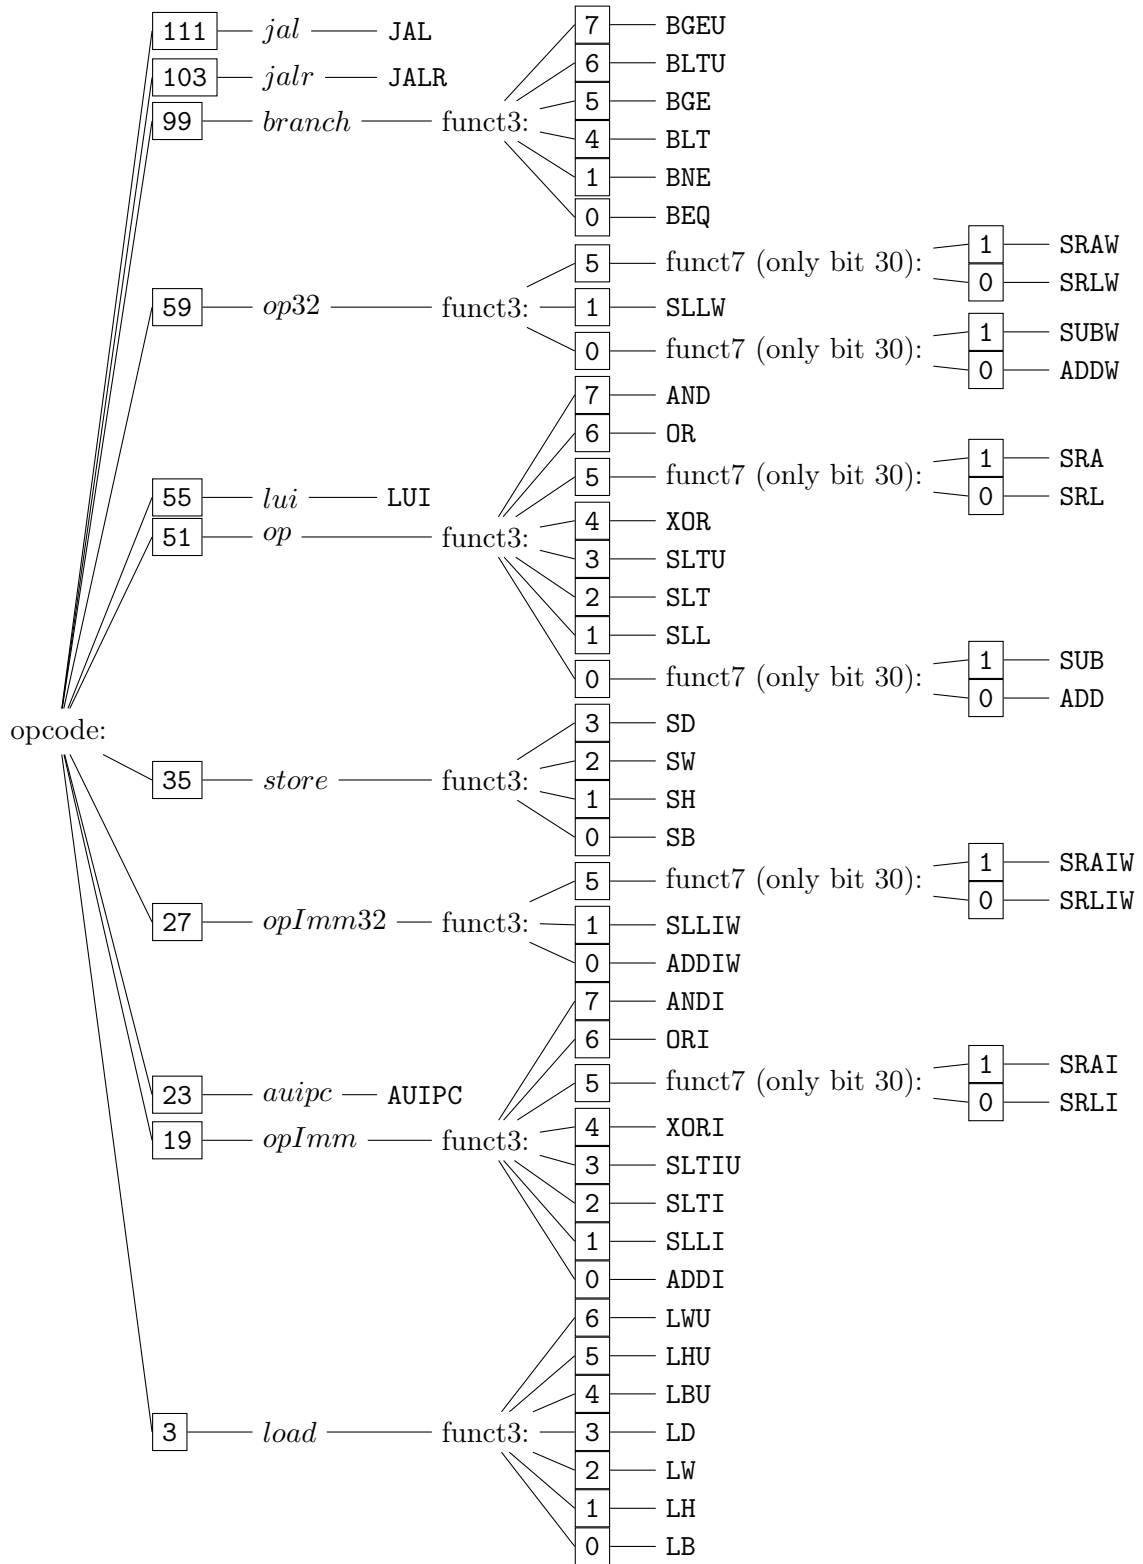
**Figure 3:** Decision tree to find the correct operation based on the current instruction

| | | |
|---|---|---|
| ⟨64bitHex⟩ | ::= | 16 digits of [0-9a-fA-F] |
| ⟨32bitHex⟩ | ::= | 8 digits of [0-9a-fA-F] |
| ⟨16bitHex⟩ | ::= | 4 digits of [0-9a-fA-F] |
| ⟨8bitHex⟩ | ::= | 2 digits of [0-9a-fA-F] |
| ⟨address⟩ | ::= | up to 16 digits of [0-9a-fA-F] |
| ⟨comment⟩ | ::= | printable characters without \n or :, up to a total line length of 80 |
| ⟨memContent⟩ | ::= | ⟨8bitHex⟩ |
| | | \| ⟨16bitHex⟩ |
| | | \| ⟨32bitHex⟩ |
| | | \| ⟨64bitHex⟩ |
| ⟨cell⟩ | ::= | ⟨address⟩:⟨memContent⟩ [#⟨comment⟩]\n |
| | | \| ⟨cell⟩⟨cell⟩ |
| ⟨regNum⟩ | ::= | 0 \|...\| 31 |
| ⟨reg⟩ | ::= | PC:⟨64bitHex⟩ [#⟨comment⟩]\n |
| | | \| x⟨regNum⟩:⟨64bitHex⟩ [#⟨comment⟩]\n |
| | | \| ⟨reg⟩⟨reg⟩ |
| ⟨memory⟩ | ::= | MEMORY:\n⟨cell⟩ |
| ⟨registers⟩ | ::= | REGISTERS:\n⟨reg⟩ |
| ⟨state⟩ | ::= | ⟨registers⟩\n⟨memory⟩\n |

**Figure 4:** Construction of .state files

is little-endian, the rightmost byte is placed at the starting address. From there the next byte on the left is placed at the next higher address. Additionally, comments can be added after a #. The total length of the line should not extend over 80 characters.

# 3 BTOR2

The second foundation of my benchmarks is BTOR2, a word-level model checking format published by A. Niemetz et al. [8]. Before explaining the format, an overview of bounded model checking is necessary.

## 3.1 Bounded Model Checking

Bounded model checking (BMC) is a formal verification technique employed to detect errors in hardware or software systems by systematically exploring the state space of a finite-state model up to a specified bound, typically defined by the number of iterations or steps. As described by A. Biere in the "Handbook of Satisfiability", BMC is primarily utilized for falsification and testing, with a focus on identifying violations of temporal properties [9]. Nevertheless, BMC can also be extended to prove properties within the given bound.

In practice, BMC translates the verification problem into a satisfiability problem, determining whether a property violation can occur within the specified bound. The model comprises a finite state machine and a set of properties to be verified. The model checker systematically explores all possible state transitions up to the bound and evaluates whether the property holds. If a violation is detected, the tool generates a *witness*, which is a trace demonstrating how the property is violated. If no violation is found within the bound, the system is considered safe up to that bound, although this does not guarantee correctness for all possible executions.

For word-level hardware model checking, the BTOR2 format has become a de facto standard for describing models and is currently used in the "Hardware Model Checking Competition" [10]. Another format is AIGER [11], from which BTOR2 is derived [8], and which is used in the bit-level track of the HWMCC [10]. In software model checking, the "Competition on Software Verification" utilizes the C and Java programming languages as input formats [12]. Additionally, a translator from BTOR2 models to C programs has been presented to bridge the gap between hardware and software verification [13]. As this work focuses on BTOR2, the following section provides a detailed overview of the format.

## 3.2 The BTOR2 Language

In BTOR2, each line represents either a sort or a node, with the line number typically serving as an identifier. A sort functions similarly to a type, defining either the length of a bitvector or the size of an array of bitvectors. Nodes represent values of a defined sort and can be constants, operations, constraints, or properties. These values can be referenced by their node identifier, i.e., the line number. The syntax of BTOR2 is detailed in [8, Figure 1], and the available operators are listed in [8, Table 1].

Key features of BTOR2 include its support for sequential operations, which facilitates the implementation of a RISC-V structure. The primary feature is the `state` operator, which defines a node that is updated sequentially. An `init` node assigns an initial value to this state, while a `next` node specifies its subsequent value. `bad` nodes can be used to define endpoints for a model, indicating either the occurrence of an unintended event or, as in this work, the discovery of the intended information. In both cases, the resulting model produces a witness. Additionally, an `input` node allows an input to the model, with assignments to this node handled by the model checker. These inputs can be constrained by `constraint` nodes, which describe invariants for the inputs of the model. An example model with these sequential nodes is shown in Figure 5. This

14

| BTOR2 model | | | | | Comments |
|---|---|---|---|---|---|
| 1 | sort | bitvec | 1 | | *bit "type"* |
| 2 | sort | bitvec | 8 | | *byte "type"* |
| 3 | constd | 2 | 99 | | *comparison constant* |
| 4 | zero | 2 | | | |
| 5 | input | 2 | | *i1* | *input i1* |
| 6 | input | 2 | | *i2* | *input i2* |
| 7 | eq | 1 | 6 5 | | |
| 8 | constraint | -6 | | | *i1 ≠ i2 must hold* |
| 9 | state | 2 | | *accu* | *sequential node accu* |
| 10 | init | 2 | 9 4 | | *initialization accu* |
| 11 | add | 2 | 6 5 | | *i1 + i2* |
| 12 | next | 2 | 9 11 | | *next accu is i1 + i2* |
| 13 | eq | 1 | 3 9 | | *accu = constant* |
| 14 | bad | 13 | | | *property: accu ≠ constant* |

**Figure 5:** An example BTOR2 model finding two numbers that are not equal and added together equal 99

model can now be checked by a model checker, which should produce a witness with an assignment of the two inputs $i1$ and $i2$ such that $s1 = i1 + i2 = constant$. Let us examine this witness in more detail.

## 3.3 The BTOR2 Witness

Running the model in Figure 22 through BtorMC [8] with the option `-trace-gen-full` produces the complete witness shown in Figure 6. The syntax of BTOR2 witnesses is described in [8, Figure 2], but I will explain the example witness in Figure 6 line by line for clarity.

The witness begins with `sat`, indicating that a property in the model is satisfiable. The second line specifies the property that was violated. In this case, only one `bad` node exists in the model, and the witness shows that `b0`, meaning the first occurring `bad` node, was violated.

```
sat
b0
#0
0 00000000 accu#0
@0
0 11101000 i1@0
1 01111011 i2@0
#1
0 01100011 accu0#1
@1
0 00000100 i1@1
1 00000000 i2@1
.
```

**Figure 6:** The complete witness of the model in Figure 5

Subsequent lines list the iterations of the counterexample for the property. For each sequential iteration, the witness first presents—marked with #$x$, where $x$ is the iteration number—a representation of all states in the current iteration with their respective values in binary. Second, marked with @$x$, all inputs for the iteration are listed, similar to the states. Note that the nodes of the inputs and the state in Figure 5 include a symbol, which is used in the witness to name the states and inputs. Examining the witness, it is evident that the counterexample requires two iterations to reach the violated property. In the first iteration, the state `accu` is initialized with 0, and `i1` and `i2` are assigned values that add up to 99 (01100011b). In the next iteration, this sum is set as the new value for `accu`, which violates the property.

# 4 Transforming RISC-V to BTOR2

This chapter addresses the central problem of this thesis: transforming a RISC-V state into the BTOR2 format for benchmarking purposes. F. Schrögendorfer conducted similar work in his master's thesis "Bounded Model Checking in Lockless Programs" [14], where he describes, among other topics, an encoding concept for a minimal machine in a multiprocessor context [14, Chapter 2]. In [14, Chapter 8], he outlines a method to encode programs for his machine model into a BTOR2 model. This approach cannot be directly replicated here, as his model assumes the entire program is known at encoding time, whereas I aim to preserve the RISC-V property that allows for self-modifying programs during execution. If this property were to be disregarded, it would be possible to analyze the complete behavior of a program by parsing the memory, but this is beyond the scope of this work. Even so, I let myself inspire by his structuring of the model.

## 4.1 The Concept

To successfully execute a RISC-V instruction, three fundamental steps must occur in sequence:

- Fetch the current instruction from memory

- Identify the instruction

- Execute the instruction

Due to the fixed instruction length of RISC-V, as mentioned in Section 2.2, fetching the current instruction is straightforward. Ultimately, a node is required that retrieves a *word* from memory at the location specified by *pc*.

For basic identification, the *opcode* must be extracted and checked. Depending on the opcode, further distinctions between instructions require extracting and checking $funct3$ and, if necessary, $funct7$. Ultimately, a node for each instruction is needed, holding a boolean value indicating whether this instruction was fetched.

To execute the instruction, the values of the immediate *imm* and, if used, the registers $rs1$ and $rs2$ must be extracted. All instructions only modify $rd$, $pc$, or memory. Therefore, the next-state logic can be generalized for these three cases.

Memory is only modified when a store instruction is identified. As all store instructions share the same type, computing the memory address is consistent across them. The final step is overwriting the memory at this address.

For the $pc$, except for jump commands, it always increments to point to the next instruction. The two unconditional jumps, `JAL` and `JALR`, must be handled separately. For branch instructions, after determining whether the relevant condition for the instruction holds, a general approach can be applied, as all branch instructions execute the same operation from this point onward.

With $rd$, generalization across instructions is not feasible. However, it is possible to generalize across all possible registers by adding a check in each register's update function to determine whether the register in question is $rd$.

## 4.2 Encoding

For improved visualization in the BTOR2 code, all sort-IDs are marked in gray, all node-IDs in red, and all non-ID numbers in blue. As described in the BTOR2 syntax [8, Figure 1], each line can have an accompanying symbol. Unfortunately, these

cannot be used as aliases for the line numbers, but for clarity, in the following figures I use them as such aliases. This allows each new figure to start with the relative line number n, making it feasible to describe repetitive addition to the model with algorithms. From these, another color, green, is included. This corresponds to an algorithm variable. In contrast to the other colors, green parts of symbols must be replaced with the variables value. As a small example, if $i = 2$ and $n = 7$,

n constd W i *iConst*

shall result in

7 constd 5 2 *2Const*

being added to the model.

It is implied that n is sufficiently incremented after adding new nodes to the model so that IDs do not overlap. With this, the following sections describe how a BTOR2 model is constructed from a RISC-V state file.

### 4.2.1 Constants

First in the model are the sorts needed. These are comparable to types and are used to define the bit width of a node. After them follow a first set of constants I called "non-progressive", which includes bitmasks, shift distances and also the values of all implemented opcodes. This is shown in Figure 7.

Of note is the representation of memory as an array of addressable memory cells, each 1 byte. The chosen address space of 16 bits is significantly smaller than the expected 64-bit address space, but representing a 64-bit addressable memory with $2^{64}$ bytes ($\approx$ 18 exabytes) is not feasible. Therefore, I selected a 16-bit address space as a practical minimum, providing approximately 65kB and supporting programs with potentially over 10,000 instructions, which I consider sufficient for most use cases. The encoding is implemented so that the address space can be modified as needed.

Following this first set of constants are the "progressive" constants. To be able to e.g.

compute which registers are encoded in the instruction or which registers to change, I need a constant for each register number. This is excellent as a minimal example for an algorithm, which is shown in Algorithm 1.

**for** i from 0 to 31 **do**
    add to model:
        n   constd   W   i        $iConst$
**end**

**Algorithm 1:** progressive constants for encoding RISC-V in BTOR2

## 4.2.2 State Representation

The next logical step is defining a representation of a RISC-V state. This is straightforward, as shown in Figure 8. I also introduced a flag for each register in my code to track whether the register was written to, enabling the transformation of a witness to a state file containing only the relevant registers. As these flags do not affect the operation of the BTOR2 model and are only included for an aesthetic choice, they are not included in my description and will not be discussed further.

## 4.2.3 Initialization

To initialize a state in BTOR2 from a RISC-V state file, the values in the registers must be loaded as constants, and for each memory address mentioned in the state file, the value and address must be loaded as constants. Due to the inability to represent a full 64-bit address space, I must manage the reduction of the address space from the state file to the BTOR2 model. I chose to initialize only the addresses up to the BTOR2 model's address space maximum and omit all others from the state file, as this provides the most predictable behavior. All addresses not mentioned in the state file are zero-initialized. Finally, these constants are used to initialize the state. For the registers, this is straightforward; for memory, all memory addresses are first

```
 1   sort    bitvec   1              Bool
 2   sort    bitvec   16             AS
 3   sort    bitvec   8              B
 4   sort    bitvec   16             H
 5   sort    bitvec   32             W
 6   sort    bitvec   64             D
 7   sort    array    2          3   Mem

 8   one     Bool                    true
 9   zero    Bool                    false
10   one     AS                      addressInc
11   constd  AS       4              pcInc
12   zero    B                       emptyCell
13   one     W                       bitPicker
14   zero    D                       emptyReg

15   consth  W        01F            5Bitmask
16   consth  W        03F            6Bitmask
17   consth  W        07F            7Bitmask
18   consth  W        0FFF           12Bitmask
19   consth  W        0FFFFF         20Bitmask

20   constd  W        7              shiftToRd
21   constd  W        15             shiftToRs1
22   constd  W        20             shiftToRs2
23   constd  W        12             shiftToFunct3
24   constd  W        25             shiftToFunct7
25   constd  W        5              shiftBy5
26   constd  W        11             shiftBy11

27   constd  W        3              load
28   constd  W        19             opImm
29   constd  W        23             auipc
30   constd  W        27             opImm32
31   constd  W        35             store
32   constd  W        51             op
33   constd  W        55             lui
34   constd  W        59             op32
35   constd  W        99             branch
36   constd  W        103            jalr
37   constd  W        111            jal
```

**Figure 7:** Sorts and non-progressive Constants for encoding RISC-V in BTOR2

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| (n + 0) | state | D | x0 | | (n + 17) | state | D | x17 |
| (n + 1) | state | D | x1 | | (n + 18) | state | D | x18 |
| (n + 2) | state | D | x2 | | (n + 19) | state | D | x19 |
| (n + 3) | state | D | x3 | | (n + 20) | state | D | x20 |
| (n + 4) | state | D | x4 | | (n + 21) | state | D | x21 |
| (n + 5) | state | D | x5 | | (n + 22) | state | D | x22 |
| (n + 6) | state | D | x6 | | (n + 23) | state | D | x23 |
| (n + 7) | state | D | x7 | | (n + 24) | state | D | x24 |
| (n + 8) | state | D | x8 | | (n + 25) | state | D | x25 |
| (n + 9) | state | D | x9 | | (n + 26) | state | D | x26 |
| (n + 10) | state | D | x10 | | (n + 27) | state | D | x27 |
| (n + 11) | state | D | x11 | | (n + 28) | state | D | x28 |
| (n + 12) | state | D | x12 | | (n + 29) | state | D | x29 |
| (n + 13) | state | D | x13 | | (n + 30) | state | D | x30 |
| (n + 14) | state | D | x14 | | (n + 31) | state | D | x31 |
| (n + 15) | state | D | x15 | | (n + 32) | state | AS | pc |
| (n + 16) | state | D | x16 | | (n + 33) | state | Mem | memory |

**Figure 8:** State representation for encoding

written into a placeholder array, which is then used to initialize the actual memory. Due to BTOR2 constraints, these constants must be defined **before** the states, but initialization with the values must occur after the states. Thus, this initialization process **wraps around** the state representation. The generation of constants is shown in Algorithm 2 and comes **before** Figure 8, while the actual initialization is shown in Algorithm 3 and comes **after** Figure 8.

### 4.2.4 Fetching the Current Instruction

To fetch the current instruction, I read the four bytes of the instruction and concatenate them, as shown in Figure 9.

$truePc \leftarrow$ value of pc in state file
$maxPc \leftarrow$ number of addresses in BTOR2 model
pcValue $\leftarrow truePc$ modulo $maxPc$
add to model:

| n | constd | AS | pcValue | pcConst |
|---|--------|----|---------|---------|

**for** every register $x_i$ **do**
    **if** register is initialised in state file **then**
        $registerValue \leftarrow$ value of $x_i$
        **if** $registerValue \neq 0$ **then**
            add to model:

| n | constd | D | registerValue | $x_i$Const |
|---|--------|---|---------------|------------|

        **end**
    **end**
**end**

add to model:

| (n + 0) | state | Mem | | | | memPH |
|---------|-------|-----|--|--|--|-------|
| (n + 1) | init | Mem | (n + 0) | emptyCell | | |

$lastPH \leftarrow memPH$
$initialCells \leftarrow$ initialised memory cells in state file with address under maxPc
**for** every cell $c$ in $cutInitialCells$ **do**
    $address \leftarrow$ address of $c$
    $value \leftarrow$ value of $c$
    add to model:

| (n + 0) | constd | AS | address | | | |
|---------|--------|----|---------|--|--|--|
| (n + 1) | constd | B | value | | | |
| (n + 2) | write | Mem | $lastPH$ | (n + 0) | (n + 1) | PHAfterC |

    $lastPH \leftarrow PHAfterc$
**end**
keep $lastPH$ for initialisation

**Algorithm 2:** Generating initialisation constants from state file in BTOR2

**for** every register $x_i$ **do**
    **if** $x_i$Const was defined **then**
        add to model:

| n | init | D | $x_i$ | $x_i$Const |
|---|------|---|-------|------------|

    **end**
**end**

With $lastPH$ from Algorithm 2
add to model:

| n | init | Mem | memory | lastPh |
|---|------|-----|--------|--------|

**Algorithm 3:** Initialising states in the BTOR2 model

| (n + 0) | add | AS | addressInc | pc |  |
|---------|-----|-----|-----------|-----|--|
| (n + 1) | add | AS | addressInc | (n + 0) |  |
| (n + 2) | add | AS | addressInc | (n + 1) |  |
| (n + 3) | read | B | memory | pc |  |
| (n + 4) | read | B | memory | (n + 0) |  |
| (n + 5) | read | B | memory | (n + 1) |  |
| (n + 6) | read | B | memory | (n + 2) |  |
| (n + 7) | concat | H | (n + 4) | (n + 3) |  |
| (n + 8) | concat | H | (n + 6) | (n + 5) |  |
| (n + 9) | concat | W | (n + 8) | (n + 7) | instr |

**Figure 9:** Fetching the current instruction from memory

### 4.2.5 Deconstruction of the Instruction

With the instruction available, it can be deconstructed to extract the *opcode*, *rd*, *rs*1, *rs*2, *funct*3, *funct*7, and *imm*. For everything except *imm*, this can be accomplished by shifting and masking, as shown in Figure 10.

The immediate, however, must first be constructed from its subfields, which are referenced in Figure 1. In the BTOR2 model, this is shown in Figure 11. This is the same method I used to get the immediate in the RISC-V simulation Section 2.3.2 Three points are noteworthy:
First, some immediate subfields overlap exactly. This is utilized in lines (n + 1) with the overlap of $imm[11 : 5]$ for I- and S-type, and (n + 21) with J- and B-types $imm[10 : 5]$ overlap. Second, as described in Section 2.2, the immediate is always sign-extended. This is achieved using arithmetic right shifts, which perform sign extension and correctly position the highest immediate bit. Third, at line (n + 8), sign extension requires a right shift by 19. As this matches the opcode for arithmetic instructions with immediate, I reused this constant.

Now, *iTypeImm*, *sTypeImm*, *bTypeImm*, *uTypeImm*, and *jTypeImm* are available. However, it is preferable to have a single node *imm* referencing the immediate value regardless of instruction. This is accomplished in Figure 12, where booleans are defined to check all opcodes that are neither R-type nor I-type. Then, if-then-else nodes are chained to select instructions of J-type, U-type, B-type, or S-type. If the instruction is none of these, I default to I-type, as R-type does not use an immediate value. Finally, *imm* is extended to the 64-bit width required by RV64I.

At this stage, the values of the designated *rs*1 and *rs*2 registers can also be extracted. This is shown for *rs*1 in Figure 4; the process is identical for *rs*2, with only the names changed. The starting equality comparisons can be omitted for *rs*2, as they are already defined for *rs*1 and can be referenced.

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | and | W | *instr* | *7Bitmask* | *opcode* |
| (n + 1) | srl | W | *instr* | *shiftToRd* | |
| (n + 2) | and | W | (n + 1) | *5Bitmask* | *rd* |
| (n + 3) | srl | W | *instr* | *shiftToRs1* | |
| (n + 4) | and | W | (n + 3) | *5Bitmask* | *rs1* |
| (n + 5) | srl | W | *instr* | *shiftToRs2* | |
| (n + 6) | and | W | (n + 5) | *5Bitmask* | *rs2* |
| (n + 7) | srl | W | *instr* | *shiftToFunct3* | |
| (n + 8) | and | W | (n + 7) | *shiftRd* | *funct3* |
| (n + 9) | srl | W | *instr* | *shiftToFunct7* | *funct7* |

**Figure 10:** Extraction of values from the instruction without *imm*

---

**for** i from 1 to 31 **do**
    add to model:

| n | eq | Bool | *rs1* | *iConst* | *isRs1Xi* |
|---|---|---|---|---|---|

**end**
add to model:

| n | ite | D | *isRs1X1* | *x1* | *x0* | *checkX1* |
|---|---|---|---|---|---|---|

**for** i from 2 to 30 **do**
    add to model:

| n | ite | D | *isRs1Xi* | *xi* | *checkX(i - 1)* | *checkXi* |
|---|---|---|---|---|---|---|

**end**
add to model:

| n | ite | D | *isRs1X31* | *x31* | *checkX30* | *rs1val* |
|---|---|---|---|---|---|---|

**Algorithm 4:** Extracting the value of the register designated by *rs1*

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | sra | W | *instr* | *shiftToRs2* | *iTypeImm* |
| (n + 1) | and | W | *iTypeImm* | *-5Bitmask* | *s[11:5]* |
| (n + 2) | add | W | *s[11:5]* | *rd* | *sTypeImm* |
| (n + 3) | and | W | *rd* | *-bitPicker* | *b[4:0]* |
| (n + 4) | and | W | *funct7* | *6Bitmask* | |
| (n + 5) | sll | W | (n + 4) | *shiftBy5* | *b[10:5]* |
| (n + 6) | and | W | *bitPicker* | *rd* | |
| (n + 7) | sll | W | (n + 6) | *shiftBy11* | *b[11]* |
| (n + 8) | sra | W | *instr* | *mathI* | |
| (n + 9) | and | W | (n + 8) | *12Bitmask* | *b[31:12]* |
| (n + 10) | add | W | *b[10:5]* | *b[4:0]* | |
| (n + 11) | add | W | *b[11]* | (n + 10) | |
| (n + 12) | add | W | *b[31:12]* | (n + 11) | *bTypeImm* |
| (n + 13) | and | W | *instr* | *-12Bitmask* | *uTypeImm* |
| (n + 14) | and | W | *rs2* | *-bitPicker* | *j[4:0]* |
| (n + 15) | and | W | *rs2* | *bitPicker* | |
| (n + 16) | sll | W | (n + 15) | *shiftBy11* | *j[11]* |
| (n + 17) | sll | W | *funct3* | *shiftToFunct3* | *j[14:12]* |
| (n + 18) | sll | W | *rs1* | *shiftToRs1* | *j[19:15]* |
| (n + 19) | sra | W | *instr* | *shiftBy11* | |
| (n + 20) | and | W | (n + 19) | *-20Bitmask* | *j[31:20]* |
| (n + 21) | add | W | *b[10:5]* | *j[4:0]* | |
| (n + 22) | add | W | *j[11]* | (n + 21) | |
| (n + 23) | add | W | *j[14:12]* | (n + 22) | |
| (n + 24) | add | W | *j[19:15]* | (n + 23) | |
| (n + 25) | add | W | *j[31:20]* | (n + 24) | *jTypeImm* |

**Figure 11:** Extraction of all *imm* types from the instruction

| | | | | | | |
|---|---|---|---|---|---|---|
| (n + 0) | eq | Bool | *opcode* | *store* | | *isSType* |
| (n + 1) | eq | Bool | *opcode* | *branch* | | *isBType* |
| (n + 2) | eq | Bool | *opcode* | *auipc* | | |
| (n + 3) | eq | Bool | *opcode* | *lui* | | |
| (n + 4) | or | Bool | (n + 2) | (n + 3) | | *isUType* |
| (n + 5) | eq | Bool | *opcode* | *jal* | | *isJType* |
| | | | | | | |
| (n + 6) | ite | W | *isSType* | *sTypeImm* | *iTypeImm* | *checkS* |
| (n + 7) | ite | W | *isBType* | *bTypeImm* | *checkS* | *checkB* |
| (n + 8) | ite | W | *isUType* | *uTypeImm* | *checkB* | *checkU* |
| (n + 9) | ite | W | *isJType* | *jTypeImm* | *checkU* | *imm32* |
| | | | | | | |
| (n + 10) | sext | D | *imm32* | *32* | | *imm* |

**Figure 12:** Finding the correct immediate by opcode

| | | | | | |
|---|---|---|---|---|---|
| (*isJALR* already exists) | | | | | |
| n | and | Bool | *isLoad* | *is5Funct3* | *isLHU* |
| (n + 0) | consth | W | *20* | | |
| (n + 1) | eq | Bool | *funct7* | (n + 0) | |
| (n + 2) | and | Bool | *is0Funct3* | (n + 1) | |
| (n + 3) | and | Bool | *isLoad* | (n + 2) | *isSUBW* |

**Figure 13:** Instruction detection of JALR, LHU and SUBW as described in Algorithm 5

### 4.2.6 Instruction Detection

For the next-state logic, it is essential to determine the current command. Therefore, I defined a check *isInstruction* for each instruction. As this is repetitive, Algorithm 5 describes a generalized approach to obtain these booleans. An example for each instruction subgroup in Algorithm 5 is provided in Figure 13. The funct7 checks from the *needsf7* subgroup can be reused if multiple instructions share the same funct7.

### 4.2.7 Next-State Logic

The next-state logic is the core of the model. Almost everything else supports this point. The goal is to create the changes each instruction would make and then apply

add to model:

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | eq | Bool | opcode | load | isLoad |
| (n + 1) | eq | Bool | opcode | opImm | isOpImm |
| (n + 2) | eq | Bool | opcode | auipc | isAUIPC |
| (n + 3) | eq | Bool | opcode | opImm32 | isOpImm32 |
| (n + 4) | eq | Bool | opcode | store | isStore |
| (n + 5) | eq | Bool | opcode | op | isOp |
| (n + 6) | eq | Bool | opcode | lui | isLUI |
| (n + 7) | eq | Bool | opcode | op32 | isOp32 |
| (n + 8) | eq | Bool | opcode | branch | isBranch |
| (n + 9) | eq | Bool | opcode | jalr | isJALR |
| (n + 10) | eq | Bool | opcode | jal | isJAL |

**for** i from 0 to 7 **do**

    add to model:

| | | | | | |
|---|---|---|---|---|---|
| n | eq | Bool | funct3 | iConst | isiFunct3 |

**end**

*onlyOp* ← [LUI, AUIPC, JAL, JALR]

*needsf7* ← [SRL, SRA, SRLI, SRAI, SRLW, SRAW, SRLWI, SRAWI, ADD, SUB, ADDW, SUBW]

*rest* ←[ all other instructions ]

**for** all instructions I in *onlyOp* **do**

    *isI* is already defined

**end**

**for** all instructions I in *rest* **do**

    *opname* ← opcode name of I

    *f3val* ← expected funct3 of I as digit

    add to model:

| | | | | |
|---|---|---|---|---|
| n | and | Bool | isopname isf3valFunct3 | isI |

**end**

**for** all instructions I in *needsf7* **do**

    *opname* ← opcode name of I

    *f3val* ← expected funct3 of I as digit

    *f7hex* ← expected funct7 of I as hexadecimal number

    add to model:

| | | | | |
|---|---|---|---|---|
| (n + 0) | consth | W | f7hex | |
| (n + 1) | eq | Bool | funct7 | (n + 0) |
| (n + 2) | and | Bool | isf3valFunct3 | (n + 1) |
| (n + 3) | and | Bool | isopname | (n + 2) isI |

**end**

**Algorithm 5:** Generalised approach to instruction detection

only the changes specific to the instruction in the state. Each state node in the model must have an accompanying next node to function correctly. First, the changed values are computed.

## Computing new Values for each Instruction

It is unnecessary to detail all instructions, as this simply follows the RV64I ISA. Instead, I provide examples for each group of instructions as divided in Table 1. These examples are found in Figure 14, specifically for

- `AUIPC` in Figure 14.1,

- `JALR` in Figure 14.2,

- `BEQ` in Figure 14.3,

- `LHU` in Figure 14.4,

- `SD` in Figure 14.7,

- `ANDI` in Figure 14.5,

- `SLLIW` in Figure 14.8,

- `SLT` in Figure 14.6, and

- `SUBW` in Figure 14.9.

These examples contain some overlaps that can be utilized, such as for load and store instructions or the cut-to-32bit values needed for word instructions. The `SD` example demonstrates that all other store instructions are interim results of preparing `SD`. Load instructions are similar, but each requires sign extension to 64 bits. Also, between `SLLIW` and `SUBW`, only one node *rs1val32* suffices. In the generated model, this is taken into account but here I avoided it, so the examples are independent from each other. With this, we have everything necessary to define the next state.

30

| n | add | D | *imm* | *pc* | *rdAUIPC* |
|---|-----|---|-------|------|-----------|

**14.1:** AUIPC

| (n + 0) | add | AS | *pc* | *pcInc* | *nextPc* |
|---------|-----|----|------|---------|----------|
| (n + 1) | add | D | *imm* | *rs1val* | |
| (n + 2) | and | D | *-1Const* | (n + 1) | |
| (n + 3) | slice | AS | (n + 2) | *15* | *pcJALR* |
| (n + 4) | uext | D | *nextPc* | *48* | *rdJALR* |

**14.2:** JALR

| (n + 0) | add | AS | *pc* | *pcInc* | | *nextPc* |
|---------|-----|----|------|---------|---|----------|
| (n + 1) | slice | AS | *imm* | *15* | *0* | *ImmAS* |
| (n + 2) | add | AS | *pc* | *ImmAS* | | *pcBranch* |
| (n + 3) | eq | Bool | *rs1val* | *rs2val* | | |
| (n + 4) | ite | AS | (n + 3) | *pcBranch* | *nextPc* | *pcBEQ* |

**14.3:** BEQ

| (n + 0) | add | D | *rs1val* | *imm* | | |
|---------|-----|---|----------|-------|---|---|
| (n + 1) | slice | AS | (n + 0) | *15* | *0* | |
| (n + 2) | add | AS | (n + 1) | *addressInc* | | |
| (n + 3) | read | B | *memory* | (n + 1) | | |
| (n + 4) | read | B | *memory* | (n + 2) | | |
| (n + 5) | concat | H | (n + 3) | (n + 4) | | |
| (n + 6) | uext | D | (n + 5) | *48* | *0* | *rdLHU* |

**14.4:** LHU

| n | and | D | *rs1val* | *imm* | *rdANDI* |
|---|-----|---|----------|-------|----------|

**14.5:** ANDI

| (n + 0) | slt | Bool | *rs1val* | *rs2val* | |
|---------|-----|------|----------|----------|---|
| (n + 1) | uext | D | (n + 0) | *63* | *rdSLT* |

**14.6:** SLT

**Figure 14:** Instruction execution for chosen instructions

| | | | | | | |
|---|---|---|---|---|---|---|
| (n + 0) | add | D | *rs1val* | *imm* | | |
| (n + 1) | slice | AS | (n + 0) | *15* | *0* | |
| (n + 2) | add | AS | (n + 1) | *addressInc* | | |
| (n + 3) | add | AS | (n + 2) | *addressInc* | | |
| (n + 4) | add | AS | (n + 3) | *addressInc* | | |
| (n + 5) | add | AS | (n + 4) | *addressInc* | | |
| (n + 6) | add | AS | (n + 5) | *addressInc* | | |
| (n + 7) | add | AS | (n + 6) | *addressInc* | | |
| (n + 8) | add | AS | (n + 7) | *addressInc* | | |
| | | | | | | |
| (n + 9) | slice | B | *rs2val* | *7* | *0* | |
| (n + 10) | slice | B | *rs2val* | *15* | *8* | |
| (n + 11) | slice | B | *rs2val* | *23* | *16* | |
| (n + 12) | slice | B | *rs2val* | *31* | *24* | |
| (n + 13) | slice | B | *rs2val* | *39* | *32* | |
| (n + 14) | slice | B | *rs2val* | *47* | *40* | |
| (n + 15) | slice | B | *rs2val* | *55* | *48* | |
| (n + 16) | slice | B | *rs2val* | *63* | *56* | |
| | | | | | | |
| (n + 17) | write | Mem | *memory* | (n + 1) | (n + 9) | *memorySB* |
| (n + 18) | write | Mem | *memorySB* | (n + 2) | (n + 10) | *memorySH* |
| (n + 19) | write | Mem | *memorySH* | (n + 3) | (n + 11) | |
| (n + 20) | write | Mem | (n + 19) | (n + 4) | (n + 12) | *memorySW* |
| (n + 21) | write | Mem | *memorySW* | (n + 5) | (n + 13) | |
| (n + 22) | write | Mem | (n + 21) | (n + 6) | (n + 14) | |
| (n + 23) | write | Mem | (n + 22) | (n + 7) | (n + 15) | |
| (n + 24) | write | Mem | (n + 23) | (n + 8) | (n + 16) | *memorySD* |

**14.7:** SD

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | and | W | *imm32* | *5Bitmask* | |
| (n + 1) | slice | W | *rs1val* | *31* | *0* |
| (n + 2) | sll | W | (n + 1) | (n + 0) | |
| (n + 3) | sext | D | (n + 2) | *32* | *rdSLLIW* |

**14.8:** SLLIW

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | slice | W | *rs1val* | *31* | *0* |
| (n + 1) | slice | W | *rs2val* | *31* | *0* |
| (n + 2) | sub | W | (n + 0) | (n + 1) | |
| (n + 3) | sext | D | (n + 2) | *32* | *rdSUBW* |

**14.9:** SUBW

**Figure 14:** Continuation of Instruction execution for chosen instructions

| (n + 0) | ite | Mem | *isSB* | *memorySB* | *memory* |
|---------|-----|-----|--------|------------|----------|
| (n + 1) | ite | Mem | *isSH* | *memorySH* | (n + 0) |
| (n + 2) | ite | Mem | *isSW* | *memorySW* | (n + 1) |
| (n + 3) | ite | Mem | *isSD* | *memorySD* | (n + 2) |
| (n + 4) | next | Mem | *memory* | (n + 3) | |

**Figure 15:** Next-State logic for the memory array

| (n + 0) | ite | AS | *isBGEU* | *pcBGEU* | *nextPc* |
|---------|-----|-----|----------|----------|----------|
| (n + 1) | ite | AS | *isBLTU* | *pcBLTU* | (n + 0) |
| (n + 2) | ite | AS | *isBGE* | *pcBGE* | (n + 1) |
| (n + 3) | ite | AS | *isBLT* | *pcBLT* | (n + 2) |
| (n + 4) | ite | AS | *isBNE* | *pcBNE* | (n + 3) |
| (n + 5) | ite | AS | *isBEQ* | *pcBEQ* | (n + 4) |
| (n + 6) | ite | AS | *isJALR* | *pcJALR* | (n + 5) |
| (n + 7) | ite | AS | *isJAL* | *pcJAL* | (n + 6) |
| (n + 8) | next | AS | *pc* | (n + 7) | |

**Figure 16:** Next-State logic for the pc register

## The Next Memory

Defining the next memory array is straightforward. All store instructions are cascaded through if-then-else nodes, with the final 'else' set as the current memory array; if no 'if' matches, the array remains unchanged. This is shown in Figure 15.

## The Next pc

For the next pc, the approach is similar, as shown in Figure 16. The only difference is that if no 'if' matches, pc must point to the next instruction to execute. The nextPc value was already computed for the JAL and JALR instructions and is reused here. The unconditional jumps also modify the value in rd, which is handled in the next section.

**The Next rd**

At last, the remaining registers must be updated. The procedure is defined in Figure 6. With the exception of x0, this is the same for all registers. The process is similar to defining the next memory or pc, but instead of a handful of instructions, all 39 relevant instructions must be considered, as only branch and store instructions do not modify rd. For brevity, the cascade for all relevant instructions is not shown in full in Algorithm 6, but only indicated.

add to model:

| n | next | D | x0 | x0 |
|---|------|---|----|----|

**for** i from 1 to 31 **do**

add to model:

| (n + 0) | ite | D | isLUI | rdLUI | xi |
|---------|-----|---|-------|-------|------|
| (n + 1) | ite | D | isAUIPC | rdAUIPC | (n + 0) |
| ⋮ | ite | D | ⋮ | ⋮ | ⋮ |
| (n + 47) | ite | D | isSRAW | rdSRAW | (n + 46) |
| (n + 48) | eq | Bool | rd | iConst | |
| (n + 49) | ite | D | (n + 48) | (n + 47) | xi |
| (n + 50) | next | D | xi | (n + 49) | |

**end**

**Algorithm 6:** Next-state logic for all x registers

### 4.2.8 Properties

The final step is to define properties to terminate the model checker. The primary constraint is reaching a set number of iterations, as shown in Figure 17.

Additional properties are defined to check for invalid instructions. The first checks if the opcode is valid for the model. The second constraint detects if the instruction cannot be identified even when the opcode is valid, as shown in Figure 18. The constraint in Figure 19 handles instruction-address-misaligned exceptions for jump instructions.

34

| | | | | | |
|---|---|---|---|---|---|
| (n + 0) | one | D | | | |
| (n + 1) | constd | D | *nIterations* | | |
| (n + 2) | state | D | | | *counter* |
| (n + 3) | init | D | (n + 2) | *emptyReg* | |
| (n + 4) | add | D | (n + 2) | (n + 0) | |
| (n + 5) | next | D | (n + 2) | (n + 4) | |
| (n + 6) | eq | Bool | (n + 2) | (n + 1) | |
| (n + 7) | bad | | (n + 6) | | |

**Figure 17:** Terminating the model by iteration count

Other properties can be defined, such as terminating on a specific pc value or when a register reaches a specified value. An example for this can be the addition of Figure 20 to the model. The easiest way to exclude the iteration counter when adding this

## 4.3 Testing for Correctness

To test my model, I compared its results to those of my RISC-V simulator (Section 2.3).

Given a state, both the simulation and the BTOR2 model are run with the iteration maximum set to 1. The resulting BTOR2 witness cannot be directly compared to the resulting state of the simulation. Therefore, I implemented a simple converter from witness to state [15, src/restate_witness.c]. These two states can then be compared. A shell script for this purpose is provided at [15, sh_utils/compare_iterations.sh].

To generate RISC-V states, I implemented a fuzzer [15, src/state_fuzzer.c] that generates randomized states with one valid instruction at the address of pc. The fuzzer first selects an instruction to test and fills all variable parts of the instruction, such as *rd* or *imm*. All registers relevant to the instruction are then assigned random 64-bit values. A pc value is generated to ensure the instruction fits within the limited address space of the BTOR2 model. If a jump instruction is chosen, possible address misalignment is corrected and address overflow is prevented. This simplifies later

```
(n + 0)    or     Bool   isLoad      isOpImm
(n + 1)    or     Bool   isAUIPC     (n + 0)
(n + 2)    or     Bool   isOpImm32   (n + 1)
(n + 3)    or     Bool   isStore     (n + 2)
(n + 4)    or     Bool   isOp        (n + 3)
(n + 5)    or     Bool   isLUI       (n + 4)
(n + 6)    or     Bool   isOp32      (n + 5)
(n + 7)    or     Bool   isBranch    (n + 6)
(n + 8)    or     Bool   isJALR      (n + 7)
(n + 9)    or     Bool   isJAL       (n + 8)
(n + 10)   bad           -(n + 9)

(n + 11)   or     Bool   isLUI       isAUIPC
(n + 12)   or     Bool   isJAL       (n + 11)
     ⋮     or     Bool     ⋮            ⋮
(n + 58)   or     Bool   isSRAW      (n + 57)
(n + 59)   and    Bool   -(n + 9)    (n + 58)
(n + 60)   bad           (n + 59)
```

**Figure 18:** Terminating the model on unknown instructions

```
(n + 0)    zero    AS
(n + 1)    constd  AS    3
(n + 2)    and     AS    (n + 1)    pcJAL
(n + 3)    and     AS    (n + 1)    pcJALR
(n + 4)    and     AS    (n + 1)    pcBEQ
     ⋮     and     AS    (n + 1)      ⋮
(n + 9)    and     AS    (n + 1)    pcBGEU

(n + 10)   neq     Bool  (n + 0)    (n + 2)
(n + 11)   neq     Bool  (n + 0)    (n + 3)
(n + 12)   neq     Bool  (n + 0)    (n + 4)
     ⋮     neq     Bool  (n + 0)      ⋮
(n + 17)   neq     Bool  (n + 0)    (n + 9)

(n + 18)   or      Bool  (n + 18)   (n + 17)
(n + 19)   or      Bool  (n + 19)   (n + 18)
     ⋮     or      Bool    ⋮          ⋮
(n + 24)   or      Bool  (n + 23)   (n + 22)
(n + 25)   bad           (n + 24)
```

**Figure 19:** Terminating the model on misaligned addresses

| | | | | |
|---|---|---|---|---|
| (n + 0) | constd | D | *forbiddenValue* | *badValRegNum* |
| (n + 1) | eq | Bool | *xRegNum* | *isXRegNumBadVal* |
| (n + 2) | bad | | *isXRegNumBadVal* | |

**Figure 20:** Terminating the model on a specific register value

comparison of the resulting states, as correct execution of the instruction always results in the same state, despite differences between the simulation and the BTOR2 model.

With this setup, a series of tests can be conducted. For this, I implemented a shell script [15, sh_utils/test_btor2_model.sh]. As the number of tests increases, it becomes more challenging to track failed tests. To address this, I wrote a script to aggregate all failed tests into one file and add additional information such as instruction name or immediate value [15, sh_utils/diff_logger.sh].

I have executed approximately 5,000,000 tests on this model without a single failure, which leads me to conclude that my implementation is correct.

# 5 Benchmarks

With the model implemented, I was able to evaluate its performance. All benchmarks were executed on an Intel Core i5-6200U using the btormc model checker, which is distributed with the BTOR2 format [8]. Each test was run five times, and the resulting times were averaged. I also attempted to run the tests with the model checkers AVR [16] and Pono [17]. The challenges encountered with these tools are discussed in Section 5.2.

## 5.1 Tests

I devised two basic tests, each composed of four RISC-V instructions as illustrated in Figure 21. One test includes a memory operation, while the other does not, allowing for measurement of the impact of memory operations on the model's performance. The programs for these tests are intentionally similar: both feature three instructions forming a loop and one instruction serving as a "workhorse". The test names, `add` and `write_mem`, are derived from this key instruction. The program is embedded into a state, as exemplified in Figure 22, where the `add` program is configured for 256 loops. In this setup, x1 acts as a loop limiter, x2 as a loop counter, and x3 as an accumulator. The instructions are placed in the initial bytes of memory. In contrast, the memory operation test uses x2 as an address to store the first byte of a register, with x3 serving as this register. Both tests were also implemented with increasing loop counts up to 2048 to provide a reference for processing more iterations.

```
        bge x2 x1 0x10           jump out of program if x1 = x2
 add x3 x3 x2 │                   either (add counter onto x3)
             │ sb x3 0x14(x2)     or (store the first byte of x3 at counter + 0x14)
        addi x2 x2 0x1           increment counter in x2
        jalr x0 x0 0x0          jump back to address 0
```

**Figure 21:** Base test cases for the benchmarks

```
REGISTERS:
PC:0
x1:100
x2:0

MEMORY:
0:001158E3 # BGE x2 x1 0x10
4:002181B3 # ADD x3 x3 x2
8:00110113 # ADDI x2 x2 1
c:00000067 # JALR x0 x0 0
```

**Figure 22:** Example state for benchmark add_0256

Additionally, I evaluated the impact of initialized memory on runtime. For this, I introduced tests with the prefix `fullmem`, where memory addresses 0x18 to 0xfff were filled with the bit pattern '0101'. The first four words of memory were filled with the test program. The fifth word could also be filled, but I opted to leave it zero-initialized to ensure the test terminates by jumping to this address and not finding a valid instruction. This guarantees termination.

I also investigated the impact of address space size by generating models from `add_0256`, `add_1024`, `write_mem_0256`, and `write_mem_1024` with extended address space. These tests use the prefix "extaddr_x", where x denotes the maximum address length in bits.

Furthermore, I implemented the base `add` tests in a manner similar to the BTOR2 model described by F. Schrögendorfer in his master thesis [14, Chapter 8]. These tests use the prefix "nopc", as the program counter is abstracted to activation flags for each instruction. The model generation is demonstrated using `nopc_add_0256` as

40

an example. First, the necessary sorts and constants for RISC-V are defined:

```
1   sort     bitvec   1      bool
2   sort     bitvec   8      memcell
3   sort     bitvec   16     addressspace
4   sort     bitvec   64     register
5   sort     array    3 2    Memory
6   zero     1               false
7   one      1               true
8   zero     2               emptymem
9   zero     4               emptyreg
10  zero     3               pcinit
11  consth   3        4      pcinc
12  consth   3        10     instr1_pcmod
13  consth   4        100    nloops
```

With this now the registers and memory can be defined:

```
99    state   3   pc      110   state   4   x10     121   state   4   x21
100   state   4   x0      111   state   4   x11     122   state   4   x22
101   state   4   x1      112   state   4   x12     123   state   4   x23
102   state   4   x2      113   state   4   x13     124   state   4   x24
103   state   4   x3      114   state   4   x14     125   state   4   x25
104   state   4   x4      115   state   4   x15     126   state   4   x26
105   state   4   x5      116   state   4   x16     127   state   4   x27
106   state   4   x6      117   state   4   x17     128   state   4   x28
107   state   4   x7      118   state   4   x18     129   state   4   x29
108   state   4   x8      119   state   4   x19     130   state   4   x30
109   state   4   x09     120   state   4   x20     131   state   4   x31
```

Note that node IDs 14-98 are skipped. As BTOR2 requires unique but not continuous node IDs, I assigned each register node ID to equal the register number plus 100 for clarity when writing models manually.

Next, the initial memory is defined:

```
144   state    5
145   init     5   14 8
146   consth   3   0
147   consth   2   e3
148   write    5   144 146 147
149   consth   3   1
150   consth   2   58
151   write    5   148 149 150
152   consth   3   2
153   consth   2   11
154   write    5   151 152 153
155   consth   3   3
156   consth   2   00
157   write    5   154 155 156
158   consth   3   4
159   consth   2   b3
160   write    5   157 158 159
161   consth   3   5
162   consth   2   81
163   write    5   16 161 162
164   consth   3   6
165   consth   2   21
166   write    5   16 164 165
167   consth   3   7
168   consth   2   00

169   write    5   166 167 168
170   consth   3   8
171   consth   2   13
172   write    5   169 170 171
173   consth   3   9
174   consth   2   01
175   write    5   172 173 174
176   consth   3   a
177   consth   2   11
178   write    5   175 176 177
179   consth   3   b
180   consth   2   00
181   write    5   178 179 180
182   consth   3   c
183   consth   2   67
184   write    5   181 182 183
185   consth   3   d
186   consth   2   00
187   write    5   184 185 186
188   consth   3   e
189   consth   2   00
190   write    5   187 188 189
191   consth   3   f
192   consth   2   00
193   write    5   190 191 192
```

Due to BTOR2 constraints, initial memory must be defined before the intended memory state node. Node 144 serves only for memory initialization, as previously described in Section 4.2.3.

The actual memory state and instruction flags are then defined. I renamed the flags from stmt (Schrögendorfer) to instr for clarity. An exit code is unnecessary, as differentiation between model terminations is not required; only the termination time is of interest.

```
199   state   5   memory
200   state   1   instr0
201   state   1   instr1
202   state   1   instr2
203   state   1   instr3
204   state   1   endflag
```

Next up would be the inputs and constraints for these. Both are not needed for the RISC-V model because neither are threads needed as I do not model parallel processing, nor is flushing as I use a naive memory model. And without inputs, the constraints are also nonexistent. So really, next up is initialization and transitions of the states. First is pc:

```
300   init   3   99    10          pc is zero initialized
301   add    3   99    11          normal pc operation
302   add    3   99    12          instr1 jump
303   eq     1   101   102         instr1 branch condition
304   ite    3   303   302   301   IF condition THEN jump ELSE increment
305   ite    3   200   304   99    IF instr1 THEN check ELSE leave pc
306   or     1   201   202         instr2-3
307   ite    3   306   301   305   IF instr2-3 THEN increment ELSE instr1
308   ite    3   203   10    307   IF instr4 THEN jump0 ELSE try instr2-3
309   next   3   99    308
```

As the transition for pc are quite possibly the most complex I added some explanation.
x0 and x1 are skipped for now, so next registers x2 and x3:

```
314   init   4   102   9               319   init   4   103   9
315   one    4                         320   add    4   102   103
316   add    4   102   315             321   ite    4   201   320   103
317   ite    4   202   316   102       322   next   4   103   321
318   next   4   102   317
```

x2 increments by one each time the third instruction is run and x3 adds x2 every
time the second instruction is run. All other registers and the memory do not change
during execution, so I show them in one big block:

```
310   init  4 100 9        340   next 4 112 112      361   init  4 123 9
311   next  4 100 100      341   init 4 113 9        362   next  4 123 123
312   init  4 101 13       342   next 4 113 113      363   init  4 124 9
313   next  4 101 101      343   init 4 114 9        364   next  4 124 124
323   init  4 104 9        344   next 4 114 114      365   init  4 125 9
324   next  4 104 104      345   init 4 115 9        366   next  4 125 125
325   init  4 105 9        346   next 4 115 115      367   init  4 126 9
326   next  4 105 105      347   init 4 116 9        368   next  4 126 126
327   init  4 106 9        348   next 4 116 116      369   init  4 127 9
328   next  4 106 106      349   init 4 117 9        370   next  4 127 127
329   init  4 107 9        350   next 4 117 117      371   init  4 128 9
330   next  4 107 107      351   init 4 118 9        372   next  4 128 128
331   init  4 108 9        352   next 4 118 118      373   init  4 129 9
332   next  4 108 108      353   init 4 119 9        374   next  4 129 129
333   init  4 109 9        354   next 4 119 119      375   init  4 130 9
334   next  4 109 109      355   init 4 120 9        376   next  4 130 130
335   init  4 110 9        356   next 4 120 120      377   init  4 131 9
336   next  4 110 110      357   init 4 121 9        378   next  4 131 131
337   init  4 111 9        358   next 4 121 121      379   init  5 199 193
338   next  4 111 111      359   init 4 122 9        380   next  5 199 199
339   init  4 112 9        360   next 4 122 122
```

Now only the instruction flags and with them the only property are left to handle:

```
381  init  1    200  7     instr0 executes initally
382  next  1    200  203   instr0 only after instr3
383  init  1    201  6
384  and   1    200  -303  instr0 and no branch
385  next  1    201  384   instr1 only after non-branching instr0
386  init  1    202  6
387  next  1    202  201   instr2 only after instr2
388  init  1    203  6
389  next  1    203  202   instr3 only after instr3
390  init  1    204  6
391  and   1    200  303
392  next  1    204  391   endflag if instr0 branches
400  bad   204             endflag terminates
```

And with this, the model is can be run. The whole suite of add tests can be derived from this model by changing the constant value of node 13 to the appropriate loop count.

## 5.2 Results

To start I placed all iterations based benchmarks with btormc into Table 3 and the benchmarks of extended address space into Table 4. Also to get an overview of the times, I plotted the times of Table 3 in Figure 23. From Table 4 one can assume that the extension of address space does not impact the runtime of the model checker in a meaningful way. As seen in Table 5, with rising loop counts, memory operations need increasingly more time in comparison to their respective non-memory operation benchmarks. Also with rising loop counts the impact of large amounts of memory initialised is reduced, which was to be expected.

The `nopc` benchmark is significantly faster than my model. This was to be expected, as Schrögendorfer models one RISC-V program specifically whereas I model the processor and feed it different programs by initialization. In this perspective, it is not surprising that a specialised model is faster than a generalised one.

I also attempted to benchmark with the AVR and Pono model checkers, as they ranked first and second in the 2024 hardware model checking competition. Unfortunately, AVR did not terminate within 15 minutes when checking the `add_0256` benchmark.

| loops | base | | fullmem | | nopc |
| --- | --- | --- | --- | --- | --- |
| | add | writemem | add | writemem | add |
| 0256 | 2.635 | 2.877 | 9.759 | 13.344 | 0.136 |
| 0512 | 6.195 | 7.306 | 16.402 | 24.209 | 0.268 |
| 0768 | 10.802 | 13.283 | 24.093 | 36.388 | 0.414 |
| 1024 | 16.306 | 21.004 | 32.732 | 50.376 | 0.569 |
| 1280 | 23.032 | 30.200 | 42.410 | 65.746 | 0.728 |
| 1536 | 30.669 | 41.262 | 52.961 | 83.036 | 0.898 |
| 1792 | 39.463 | 53.940 | 64.598 | 101.475 | 1.075 |
| 2048 | 48.944 | 68.521 | 77.084 | 122.189 | 1.276 |

**Table 3:** Times of iterations based benchmarks

| pc width | add_0256 | add_1024 | writemem_0256 | writemem_1024 |
| --- | --- | --- | --- | --- |
| 16 | 2.635 | 16.306 | 2.877 | 21.004 |
| 17 | 2.632 | 16.464 | 2.88 | 21.131 |
| 18 | 2.626 | 16.511 | 2.890 | 21.215 |
| 19 | 2.626 | 16.452 | 2.889 | 21.181 |
| 20 | 2.624 | 16.460 | 2.890 | 21.163 |

**Table 4:** Times of extended address space benchmarks
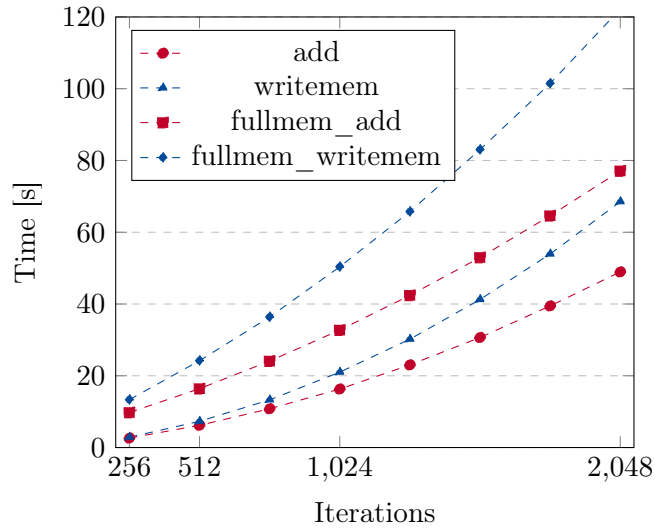


**Figure 23:** Table 3 plotted

| loops | $\dfrac{add}{writemem}$ | $\dfrac{fullmem\_add}{fullmem\_writemem}$ | $\dfrac{add}{fullmem\_add}$ | $\dfrac{writemem}{fullmem\_writemem}$ | $\dfrac{add}{nopc\_add}$ |
|---|---|---|---|---|---|
| 0256 | 0.92 | 0.73 | 0.27 | 0.22 | 19.38 |
| 0512 | 0.85 | 0.68 | 0.38 | 0.3 | 23.12 |
| 0768 | 0.81 | 0.66 | 0.45 | 0.37 | 26.09 |
| 1024 | 0.78 | 0.65 | 0.5 | 0.42 | 28.66 |
| 1280 | 0.76 | 0.65 | 0.54 | 0.46 | 31.64 |
| 1536 | 0.74 | 0.64 | 0.58 | 0.5 | 34.15 |
| 1792 | 0.73 | 0.64 | 0.61 | 0.53 | 36.71 |
| 2048 | 0.71 | 0.63 | 0.63 | 0.56 | 38.36 |

**Table 5:** Relative runtime of benchmarks

Pono determined that this benchmark is satisfiable, but required nearly six minutes using the `-smt-solver cvc5` option. With `-smt-solver bzla`, Pono was slower but functional; other solvers reported inability to handle arrays. These results suggest that model checkers considered superior in general are not necessarily optimal for my model, whereas btormc, despite being unmaintained since August 2024, performed best. I suspect that newer model checkers are optimized for handling inputs, as most competition benchmarks involve at least one input beyond a clock signal, while my model operates solely with known constants and its execution time is dependent on rapid state iteration.

Another issue with AVR and Pono is their inability to generate a "complete" witness, as produced by the btormc option `-trace-gen-full`, which provides the values of all states in the final frame. Without this feature, it is not possible to reconstruct a state from their output after execution.

# 6 Conclusion

I developed tools to transition from a state to a model and from a witness of this model back to a state. Additionally, I implemented a fuzzer for the states to verify the correct functioning of the model, as well as a set of basic tests to benchmark its performance. Finally, I presented and discussed the results of these benchmarks, concluding that model checkers which appear superior in general are worse at this specific task of mainly running iterations.

**(TODO: Check names of chapters and figs and tables and algs)**
**(TODO: version numbers for repo sources)**
**(TODO: remove todo list :))**

# ToDo Counters

To Dos: 4;    1, 2, 3, 4

Parts to extend: 0;

Draft parts: 0;

# Bibliography

[1] C. of the European Union, "Regulation (eu) 2023/1781," https://eur-lex.europa.eu/eli/reg/2023/1781/oj, 2023.

[2] A. Waterman, Y. Lee, D. A. Patterson, and K. Asanović, "The risc-v instruction set manual, volume i: Base user-level isa," UC Berkeley, Tech. Rep. UCB/EECS-2011-62, May 2011. [Online]. Available: http://www2.eecs.berkeley.edu/Pubs/TechRpts/2011/EECS-2011-62.html

[3] "History of RISC-V," https://riscv.org/about/, accessed: 15.08.2025.

[4] *The RISC-V Instruction Set Manual Volume I: Unprivileged ISA*, 2025, version 20250508. [Online]. Available: https://lf-riscv.atlassian.net/wiki/spaces/HOME/pages/16154769/RISC-V+Technical+Specifications

[5] "RISC-V-Simulator," https://github.com/Kr1mo/Risc-V-Simulator.

[6] F. Bellard, "Qemu, a fast and portable dynamic translator," in *Proceedings of the Annual Conference on USENIX Annual Technical Conference*, ser. ATEC '05. USA: USENIX Association, 2005, p. 41.

[7] J. Dupák, P. Píša, M. Štepanovský, and K. Kočí, "Qtrvsim - risc-v simulator for computer architectures classes," in *embedded world Conference 2022*. WEKA FACHMEDIEN GmbH, 2022, pp. 775–778.

[8] A. Niemetz, M. Preiner, C. Wolf, and A. Biere, "Btor2 , BtorMC and Boolector 3.0," in *Computer Aided Verification*, H. Chockler and G. Weissenbacher, Eds. Cham: Springer International Publishing, 2018, pp. 587–595.

[9] A. Biere, "Bounded model checking," in *Handbook of Satisfiability*, 2nd ed., ser. Frontiers in Artificial Intelligence and Applications, A. Biere, M. Heule, H. van Maaren, and T. Walsh, Eds. IOS Press, 2021, vol. 336, pp. 739–764. [Online]. Available: https://doi.org/10.3233/FAIA201002

[10] "HWMCC 2024," https://hwmcc.github.io/2024/.

[11] A. Biere, "The AIGER And-Inverter Graph (AIG) format version 20071012," Institute for Formal Models and Verification, Johannes Kepler University, Altenbergerstr. 69, 4040 Linz, Austria, Tech. Rep. 07/1, 2007.

[12] D. Beyer and J. Strejček, "Improvements in software verification and witness validation: SV-COMP 2025," in *Proc. TACAS (3)*, ser. LNCS 15698. Springer, 2025, pp. 151–186.

[13] D. Beyer, P.-C. Chien, and N.-Z. Lee, "Bridging hardware and software analysis with btor2c: A word-level-circuit-to-c translator," in *Tools and Algorithms for the Construction and Analysis of Systems*, S. Sankaranarayanan and N. Sharygina, Eds. Cham: Springer Nature Switzerland, 2023, pp. 152–172.

[14] F. Schrögendorfer, "Bounded Model Checking of Lockless Programs," Master's thesis, Johannes Kepler University Linz, August 2021. [Online]. Available: https://epub.jku.at/obvulihs/download/pdf/6579523

[15] "RISC-V_to_BTOR2," https://github.com/Kr1mo/RISC-V_to_BTOR2.

[16] A. Goel and K. Sakallah, "Avr: Abstractly verifying reachability," in *Tools and Algorithms for the Construction and Analysis of Systems*, A. Biere and D. Parker, Eds. Cham: Springer International Publishing, 2020, pp. 413–422.

[17] M. Mann, A. Irfan, F. Lonsing, Y. Yang, H. Zhang, K. Brown, A. Gupta, and C. Barrett, "Pono: A flexible and extensible smt-based model checker," in *Computer Aided Verification: 33rd International Conference, CAV 2021, Virtual Event, July 20–23, 2021, Proceedings, Part II.* Berlin, Heidelberg: Springer-Verlag, 2021, p. 461–474. [Online]. Available: https://doi.org/10.1007/978-3-030-81688-9_22