

Bachelor Thesis

Benchmark of RISC-V in BTOR2

Jan Krister Möller

Examiner: Dr. Mathias Fleury

University of Freiburg
Faculty of Engineering
Department of Computer Science
Chair of Computer Architecture

September 21, 2025

Writing Period

24. 06. 2025 – 24. 09. 2025

Examiner

Dr. Mathias Fleury

Declaration

I hereby declare that I am the sole author and composer of my thesis and that no other sources or learning aids, other than those listed, have been used. Furthermore, I declare that I have acknowledged the work of others by providing detailed references of said work.

I hereby also declare that my Thesis has not been prepared for another examination or assignment, either wholly or excerpts thereof.

Place, Date

Signature

Declaration on Usage of Generative AI

I hereby declare that, with the approval of my examiner, I have employed the generative AI tool “GitHub Copilot” during the preparation of this thesis solely for spellchecking and enhancing the formality of my written expression. Furthermore, I expressly confirm that this tool was not used to generate data or as a source of factual information or content for this thesis.

Place, Date

Signature

Abstract

foo bar [1] [2] [3]

Contents

1	Introduction	1
2	RISC-V	3
2.1	Overview	3
2.2	The RV64I ISA	4
2.3	Simulation of RISC-V	8
2.3.1	Representing the State of a RISC-V Processor	8
2.3.2	Running an Instruction	9
2.3.3	Saving the State of a RISC-V Processor	10
3	BTOR2	11
3.1	Bounded Model Checking	11
3.2	The BTOR2 Language	12
3.3	The BTOR2 Witness	13
4	Transforming RISC-V to BTOR2	15
4.1	The Concept	15
4.2	Encoding	17
4.2.1	Constants	17
4.2.2	State Representation	19
4.2.3	Initialization	20
4.2.4	Fetching the Current Instruction	20
4.2.5	Deconstruction of the Instruction	20

4.2.6	Instruction Detection	26
4.2.7	Next-State Logic	26
4.2.8	Constraints	31
4.3	Testing for Correctness	34
5	Benchmarks	35
5.1	Tests	35
5.2	Results	40
6	Conclusion	43
	Bibliography	48

List of Figures

1	RV64I encoding formats	5
2	State representation of RISC-V in c	9
3	Construction of .state files	10
4	An example BTOR2 model	13
5	An example BTOR2 witness	14
6	Sorts and non-progressive Constants	18
7	State representation for transforming RISC-V to BTOR2	19
8	Fetching instruction	22
9	Extraction (without immediate)	23
10	Extraction of immediate types	24
11	Finding the correct immediate	25
12	Examples for instruction detection	26
13	Examples for instruction execution	28
	13.1 AUIPC	28
	13.2 JALR	28
	13.3 BEQ	28
	13.4 LHU	29
	13.5 SD	29
	13.6 ANDI	30
	13.7 SLLIW	30
	13.8 SLT	30

	13.9 SUBW	30
14	Next-State logic for memory	30
15	Next-State logic for pc	31
16	Iterations counter constraint	32
17	Unknown Instruction constraints	33
18	Instruction address misaligned constraint	33
19	Base test cases	36
20	State of benchmark add_0256	36

List of Tables

1	RV64I Instruction Subset	7
2	Behavior of RV64I operations	8
3	Times of iterations based benchmarks	41
4	Times of extended address space benchmarks	41
5	Relative runtime of benchmarks	41

List of Algorithms

1	progressive constants	19
2	Generating initialisation constants	21
3	Initialising states	22
4	Value extraction from <i>rs1</i>	25
5	Instruction detection	27
6	Next-state logic for rd	32

1 Introduction

(TODO: Write this ?!)

2 RISC-V

As the first foundation for my benchmarks and, consequently, this thesis, I will discuss RISC-V and its operational principles.

2.1 Overview

RISC-V is an open-source instruction set architecture (ISA) first published in May 2011 by A. Waterman et al. [4]. As indicated by its name, it is based on the RISC design philosophy. **(TODO: Explain RISC (compare wiki))** Since 2015, the development of RISC-V has been coordinated by the RISC-V International Association, a non-profit corporation based in Switzerland since 2020 [5]. Its objectives include providing an *open* ISA that is freely available to all, a *real* ISA suitable for native hardware implementation, and an ISA divided into a *small* base integer ISA usable independently, for example in educational contexts, with optional standard extensions to support general-purpose software development [1, Chapter 1].

Currently, RISC-V comprises four base ISAs: RV32I, RV64I, RV32E, and RV64E, which can be extended with one or more of the 47 ratified extension ISAs [1, Preface].

(EXTEND: Additional content may be required here)

For the purposes of this work, I focus on a subset of the RV64I ISA.

2.2 The RV64I ISA

RV64I is not overly complex, but its structure is essential for understanding the subsequent work presented in this thesis. Therefore, I will explain all elements relevant to my work.

RV64I features 32 64-bit registers, labeled $x0$ – $x31$, where $x0$ is hardwired to zero across all bits. Registers $x1$ – $x31$ are general-purpose and may be interpreted by various instructions as collections of booleans, two’s complement signed binary integers, or unsigned integers. Additionally, there is a non-accessible register called pc , which serves as the program counter and holds the address of the current instruction [1, Chapters 4.1, 2.1].

In RV64I, memory addresses are 64 bits in size. As the memory model is defined to be single-byte addressable, the address space of RV64I encompasses 2^{64} bytes [1, Chapter 1.4]. The format of the memory is little endian, so the lower bits of a number are placed at lower addresses.

Like nearly all standard ISAs of RISC-V, RV64I employs a standard instruction encoding length of 32 bits, or one *word*. Only the compressed extension named C introduces instructions with a length of 16 bits [1, Chapter 1.5], but this special case is not considered here. All RV64I instructions are encoded in one of the six formats illustrated in Figure 1. These formats may consist of

- The *opcode*:

The opcode is used to differentiate between groups of instructions. It also defines the format type of the instruction.

- *rd*:

This is the destination register.



Figure 1: RV64I encoding formats, used in [1, Chapter 2.3]

- *funct3*:

This is used to differentiate between instructions with the same *opcode*.

- *rs1* & *rs2*:

These are the source registers.

- *funct7*:

This is used for further distinctions between instructions if there are more than eight instructions in an opcode group and *funct3* does not suffice.

- *imm*:

This is an immediate value. In square brackets after *imm* is designated a subfield of the immediate which is represented by these bits. From these subfields, non-defined lower bits are filled with zeros whereas the highest defined bit is sign-extended to fill all non-defined higher bits.

The design of these formats results in the following features:

- Due to RISC-V's little-endian nature, the *opcode*, which encodes the general instruction, is always read first. Further specification of the instruction via *funct3* and *funct7* is consistently located at the same positions.

- If utilized by the instruction, rd , $rs1$, and $rs2$ are also always found in the same locations, simplifying decoding.
- The highest bit of imm is always bit 31, making it straightforward to sign-extend the immediate value.

The instructions relevant to my work are listed in Table 1 I have divided the instructions in Table 1 into nine groups based on their operations.

LUI and **AUIPC** move a high immediate into rd . In the case of **AUIPC**, the pc is added to this value. **JAL** and **JALR** instructions are unconditional jumps, where for **JAL** imm is added to pc and for **JALR** imm is added to $rs1$ and set as pc . Both link to the next instruction (current $pc + 4$) in rd .

branch instructions are conditional jumps. $rs1$ is compared to $rs2$ and if the comparison holds, imm is added to pc . The comparisons are $=$ for **BEQ**, \neq for **BNE**, $<$ for **BLT**, and \geq for **BGE**. In these instructions, the values in $rs1$ and $rs2$ are handled as two's complement integers. The suffix ***U** in an instruction generally designates an unsigned operation. In this case, the values in $rs1$ and $rs2$ are handled as unsigned integers. Apart from this, they work as their counterpart without the suffix.

load instructions load values from memory at address $(rs1 + imm)$ into rd , either at Byte, Halfword, Word, or Doubleword length. By default, the value is sign-extended, and the suffix ***U** designates the loading of a non-sign-extended value. Conversely, *store* instructions write values from $rs2$ at the address $(rs1 + imm)$ to memory. Here also the distinction between the different lengths is made, and the lowest byte, halfword, word, or the whole doubleword is stored at the address.

All further instructions can be seen as generic operations, differentiated by their suffixes. To simplify the explanation process, all operations without any suffix and their behavior are listed in Table 2. This is almost exactly the group with opcode op , except the **SLTU** instruction, which is not suffix-free. However, as with all other

Instr	opcode	Type	Instr	opcode	Type	Instr	opcode	Type
LUI	<i>lui</i>	U	SB	<i>store</i>	S	ADD	<i>op</i>	R
AUIPC	<i>auipc</i>		SH			SUB		
JAL	<i>jal</i>	J	SW			SLT		
JALR	<i>jalr</i>	I	SD			SLTU		
BEQ	<i>branch</i>	B	ADDI	<i>op-imm</i>	I	XOR		
BNE			SLTI			OR		
BLT			SLTIU			AND		
BGE			XORI			SLL		
BLTU			ORI			SRL		
BGEU			ANDI			SRA		
LB			SLLI		I*	ADDW	<i>op-32</i>	R
LH	<i>load</i>	I	SRLI			SUBW		
LW			SRAI			SLLW		
LD			ADDIW	<i>op-imm-32</i>	I	SRLW		
LBU			SLLIW		I**	SRAW		
LHU			SRLIW					
LWU			SRAIW					

Table 1: Subset of RV64I instructions (**TODO: Maybe rework, not happy yet**)

instructions with the unsigned suffix, it behaves as its signed counterpart except for handling both *rs1* and *rs2* as unsigned integers.

These operations can be extended by the *I suffix, which is designated by the opcode *op-imm*. This replaces *rs2* with *imm* in the behavior. Again, SLTI can be extended to an unsigned version SLTIU, which behaves as expected. A SUBI instruction does not exist as it is redundant; its behavior can be achieved by using ADDI with a negative immediate.

Additionally, the operations ADD, SUB, SLL, SRL, and SRA can be extended with the *W suffix. This forms the group with the opcode *op-32*. In contrast to the base instructions, these new ones behave as if the registers are only 32 bits. The result is placed in the low 32 bits of *rd* and sign-extended to the full 64 bits. Overflows are ignored.

The last group is the combination of both suffixes *IW with the opcode *op-imm-32*. The behavior differs from the base instructions, as expected, by a replacement of *rs2*

Instr	Behavior
ADD	$rd := rs1 + rs2$
SUB	$rd := rs1 - rs2$
SLT	$rd := 1$ if $rs1 < rs2$ else $rd := 0$
XOR	$rd := rs1 \oplus rs2$, bitwise
OR	$rd := rs1 \vee rs2$, bitwise
AND	$rd := rs1 \wedge rs2$, bitwise
SLL	$rd := rs1$ shifted left by $rs2$, new bits are zeros
SRL	$rd := rs1$ shifted right by $rs2$, new bits are zeros
SRA	$rd := rs1$ shifted right by $rs2$, sign extend

Table 2: All suffix free operations in RV64I and their behavior. All values are handled either bitwise or as signed twos complement integers

with *imm* and only operating on 32 bits. Again, a SUBIW instruction is redundant as a negative immediate with ADDIW achieves the same result.

Compared to the full RV64I ISA, I have omitted the FENCE, ECALL, and EBREAK instructions, as without I/O interaction or an environment such as an OS or a debugger, these are not required.

2.3 Simulation of RISC-V

In a previous project (**TODO: Must be more precise**), I implemented a simulation of this subset of RV64I in C [6]. As I will use it to test my BTOR2 model, I explain the inner workings of the program here.

(TODO: related work in simulation)

First, I implemented a structure to represent a simple RISC-V processor.

2.3.1 Representing the State of a RISC-V Processor

The state requires a representation for all registers. *pc* is defined as a 64-bit integer, and the other 32 registers are implemented as an array, allowing each register to


```

typedef struct memory_cell
{
    uint64_t address;
    uint8_t content;
    struct memory_cell *next_cell;
} memory_cell;

typedef struct memory_table
{
    memory_cell *memory[TABLESIZE];
    uint64_t initialised_cells;
} memory_table;

typedef struct state
{
    uint64_t pc;
    uint64_t regs_values[32];
    bool regs_init[32];
    memory_table *memory;
} state;

```

Figure 2: State representation of a RISC-V processor in the simulation [6]

be referenced by its number. Additionally, I implemented an array of flags, one for each register, to differentiate between initialized and non-initialized registers. The memory is built from single memory cells, each holding an address and its byte of content. These are accumulated in a hash table called “memorytable”, hashing on the address. If adding a new cell causes a collision, it is appended to the cells already in the bucket, forming a linked list. These structures are shown in Figure 2.

2.3.2 Running an Instruction

After fetching the current instruction from the hash table, it must be decoded. The easiest way to do this is using a decision tree. First, I mask out the opcode and match it over all implemented opcodes. From there, either this is an endpoint and the instruction is identified, or *funct3* must be masked and matched. A final differentiation over *funct7* might be needed, but after this every leaf in the tree coincides with an instruction. At every leaf, the state can be changed according to the corresponding instruction. **(TODO: Graph of the decision tree?)**

$\langle 64bitHex \rangle$::=	up to 16 digits of [0-9a-fA-F]
$\langle 32bitHex \rangle$::=	up to 8 digits of [0-9a-fA-F]
$\langle 16bitHex \rangle$::=	up to 4 digits of [0-9a-fA-F]
$\langle 8bitHex \rangle$::=	up to 2 digits of [0-9a-fA-F]
$\langle memContent \rangle$::=	$\langle 8bitHex \rangle$ $\langle 16bitHex \rangle$ $\langle 32bitHex \rangle$ $\langle 64bitHex \rangle$
$\langle cell \rangle$::=	$\langle 64bitHex \rangle : \langle memContent \rangle$ $\langle cell \rangle \langle cell \rangle$
$\langle regNum \rangle$::=	0 ... 31
$\langle reg \rangle$::=	PC: $\langle 64bitHex \rangle \backslash n$ x $\langle regNum \rangle$: $\langle 64bitHex \rangle \backslash n$ $\langle reg \rangle \langle reg \rangle$
$\langle memory \rangle$::=	MEMORY: $\backslash n \langle cell \rangle$
$\langle registers \rangle$::=	REGISTERS: $\backslash n \langle reg \rangle$
$\langle state \rangle$::=	$\langle registers \rangle \backslash n \langle memory \rangle \backslash n$

Figure 3: Construction of .state files

2.3.3 Saving the State of a RISC-V Processor

To preserve the current state of a RISC-V processor, both the registers and memory must be stored. For this purpose, I have devised the format shown in Figure 3. The RISC-V simulation uses this format as input and output. The minimal file consists only of the two designators “REGISTERS:” and “MEMORY:” and one empty line between them. Under “REGISTERS:”, all registers can be listed with their corresponding value. Of course, x0 cannot be different from 0. I included the option to reference it nonetheless to have the complete state included. Under “MEMORY:”, after giving an address, the memory can be filled with 1-, 2-, 4-, or 8-byte sized memory content. The given address is the starting address of the content, with every byte after the first filling the next higher address.

3 BTOR2

The second foundation of my benchmarks is BTOR2, a word-level model checking format published by A. Niemetz et al. [2]. Before explaining the format, an overview of bounded model checking is necessary.

3.1 Bounded Model Checking

Bounded model checking (BMC) is a formal verification technique employed to detect errors in hardware or software systems by systematically exploring the state space of a finite-state model up to a specified bound, typically defined by the number of iterations or steps. As described by A. Biere in the “Handbook of Satisfiability”, BMC is primarily utilized for falsification and testing, with a focus on identifying violations of temporal properties [7]. Nevertheless, BMC can also be extended to prove properties within the given bound.

In practice, BMC translates the verification problem into a satisfiability problem, determining whether a property violation can occur within the specified bound. The model comprises a finite state machine and a set of properties to be verified. The model checker systematically explores all possible state transitions up to the bound and evaluates whether the property holds. If a violation is detected, the tool generates a *witness*, which is a trace demonstrating how the property is violated. If no violation is found within the bound, the system is considered safe up to that bound, although this does not guarantee correctness for all possible executions.

For word-level hardware model checking, the BTOR2 format has become a de facto standard for describing models and is currently used in the “Hardware Model Checking Competition” [8]. Another format is AIGER [9], from which BTOR2 is derived [2], and which is used in the bit-level track of the HWMCC [8]. In software model checking, the “Competition on Software Verification” utilizes the C and Java programming languages as input formats [10]. Additionally, a translator from BTOR2 models to C programs has been presented to bridge the gap between hardware and software verification [11]. As this work focuses on BTOR2, the following section provides a detailed overview of the format.

3.2 The BTOR2 Language

In BTOR2, each line represents either a sort or a node, with the line number typically serving as an identifier. A sort functions similarly to a type, defining either the length of a bitvector or the size of an array of bitvectors. Nodes represent values of a defined sort and can be constants, operations, or constraints. These values can be referenced by their node identifier, i.e., the line number. The syntax of BTOR2 is detailed in [2, Figure 1], and the available operators are listed in [2, Table 1].

Key features of BTOR2 include its support for sequential operations, which facilitates the implementation of a RISC-V structure. The primary feature is the **state** operator, which defines a node that is updated sequentially. An **init** node assigns an initial value to this state, while a **next** node specifies its subsequent value. **bad** nodes can be used to define endpoints for a model, indicating either the occurrence of an unintended event or, as in this work, the discovery of the intended information. In both cases, the resulting model produces a witness. Additionally, an **input** node allows an input to the model, with assignments to this node handled by the model checker. These inputs can be constrained by **constraint** nodes, which describe invariants for the inputs of the model. An example model with these sequential nodes is shown in Figure 4. This

BTOR2 model				Comments
1	sort	bitvec	1	<i>bit "type"</i>
2	sort	bitvec	8	<i>byte "type"</i>
3	constd	2	99	<i>comparison constant</i>
4	zero	2		
5	input	2		<i>input i1</i>
6	input	2		<i>input i2</i>
7	eq	1	6 5	
8	constraint	-6		<i>i1 ≠ i2 must hold</i>
9	state	2		<i>sequential node accu</i>
10	init	2	9 4	<i>initialization accu</i>
11	add	2	6 5	<i>i1 + i2</i>
12	next	2	9 11	<i>next accu is i1 + i2</i>
13	eq	1	3 9	<i>accu = constant</i>
14	bad	13		<i>property: accu ≠ constant</i>

Figure 4: An example BTOR2 model finding two numbers that are not equal and added together equal 99

model can now be checked by a model checker, which should produce a witness with an assignment of the two inputs $i1$ and $i2$ such that $s1 = i1 + i2 = \text{constant}$. Let us examine this witness in more detail.

3.3 The BTOR2 Witness

Running the model in Figure 20 through BtorMC [2] with the option `-trace-gen-full` produces the complete witness shown in Figure 5. The syntax of BTOR2 witnesses is described in [2, Figure 2], but I will explain the example witness in Figure 5 line by line for clarity.

The witness begins with **sat**, indicating that a property in the model is satisfiable. The second line specifies the constraint that was triggered. In this case, only one **bad** node exists in the model, and the witness shows that **b0**, meaning the first occurring **bad** node, was violated.

```

sat
b0
#0
0 00000000 accu#0
@0
0 11101000 i1@0
1 01111011 i2@0
#1
0 01100011 accu0#1
@1
0 00000100 i1@1
1 00000000 i2@1
.

```

Figure 5: The complete witness of the model in Figure 4

Subsequent lines list the iterations of the counterexample for the property. For each sequential iteration, the witness first presents—marked with $\#x$, where x is the iteration number—a representation of all states in the current iteration with their respective values in binary. Second, marked with $@x$, all inputs for the iteration are listed, similar to the states. Note that the nodes of the inputs and the state in Figure 4 include a symbol, which is used in the witness to name the states and inputs. Examining the witness, it is evident that the counterexample requires two iterations to reach the violated property. In the first iteration, the state `accu` is initialized with 0, and `i1` and `i2` are assigned values that add up to 99 (01100011b). In the next iteration, this sum is set as the new value for `accu`, which violates the property.

4 Transforming RISC-V to BTOR2

(TODO: Explain naming conventions for the model nodes)

This chapter addresses the central problem of this thesis: transforming a RISC-V state into the BTOR2 format for benchmarking purposes. F. Schrögenderfer conducted similar work in his master’s thesis “Bounded Model Checking in Lockless Programs” [3], where he describes, among other topics, an encoding concept for a minimal machine in a multiprocessor context [3, Chapter 2]. In [3, Chapter 8], he outlines a method to encode programs for his machine model into a BTOR2 model. This approach cannot be directly replicated here, as his model assumes the entire program is known at encoding time, whereas I aim to preserve the RISC-V property that allows for self-modifying programs during execution. If this property were to be disregarded, it would be possible to analyze the complete behavior of a program by parsing the memory, but this is beyond the scope of this work. Even so, I let myself inspire by his structuring of the model.

4.1 The Concept

To successfully execute a RISC-V instruction, three fundamental steps must occur in sequence:

- Fetch the current instruction from memory
- Identify the instruction

- Execute the instruction

Due to the fixed instruction length of RISC-V, as mentioned in Section 2.2, fetching the current instruction is straightforward. Ultimately, a node is required that retrieves a *word* from memory at the location specified by *pc*.

For basic identification, the *opcode* must be extracted and checked. Depending on the opcode, further distinctions between instructions require extracting and checking *funct3* and, if necessary, *funct7*. Ultimately, a node for each instruction is needed, holding a boolean value indicating whether this instruction was fetched.

To execute the instruction, the values of the immediate *imm* and, if used, the registers *rs1* and *rs2* must be extracted. All instructions only modify *rd*, *pc*, or memory. Therefore, the next-state logic can be generalized for these three cases.

Memory is only modified when a store instruction is identified. As all store instructions share the same type, computing the memory address is consistent across them. The final step is overwriting the memory at this address.

For the *pc*, except for jump commands, it always increments to point to the next instruction. The two unconditional jumps, JAL and JALR, must be handled separately. For branch instructions, after determining whether the relevant condition for the instruction holds, a general approach can be applied, as all branch instructions execute the same operation from this point onward.

With *rd*, generalization across instructions is not feasible. However, it is possible to generalize across all possible registers by adding a check in each register's update function to determine whether the register in question is *rd*.

4.2 Encoding

For improved visualization in the BTOR2 code, all sort-IDs are marked in gray, all node-IDs in red, and all non-ID numbers in blue. As described in the BTOR2 syntax [2, Figure 1], each line can have an accompanying symbol. Unfortunately, these cannot be used as aliases for the line numbers, but for clarity, in the following figures I use them as such aliases. This allows each new figure to start with the relative line number n , making it feasible to describe processes with algorithms. It is implied that n is sufficiently incremented after adding to the model so that IDs do not overlap. The following sections describe how a BTOR2 model is constructed from a RISC-V state file.

4.2.1 Constants

First, I added the sorts and non-progressive constants needed in the BTOR2 model, as shown in Figure 6. This is extended by a set of progressive constants used for comparison, e.g., against the register number. Algorithm 1 describes how these are added.

Of note is the representation of memory as an array of addressable memory cells, each 1 byte. The chosen address space of 16 bits is significantly smaller than the expected 64-bit address space, but representing a 64-bit addressable memory with 2^{64} bytes (≈ 18 exabytes) is not feasible. Therefore, I selected a 16-bit address space as a practical minimum, providing approximately 65kB and supporting programs with potentially over 10,000 instructions, which I consider sufficient for most use cases. The encoding is implemented so that the address space can be modified as needed.

(TODO: Change code to make address space modifiable by an option?)

(TODO: Explain progressive constants)

1	sort	bitvec	1		<i>Bool</i>
2	sort	bitvec	16		<i>AS</i>
3	sort	bitvec	8		<i>B</i>
4	sort	bitvec	16		<i>H</i>
5	sort	bitvec	32		<i>W</i>
6	sort	bitvec	64		<i>D</i>
7	sort	array	2	3	<i>Mem</i>
8	one	Bool			<i>true</i>
9	zero	Bool			<i>false</i>
10	one	AS			<i>addressInc</i>
11	constd	AS	4		<i>pcInc</i>
12	zero	B			<i>emptyCell</i>
13	one	W			<i>bitPicker</i>
14	zero	D			<i>emptyReg</i>
15	consth	W	01F		<i>5Bitmask</i>
16	consth	W	03F		<i>6Bitmask</i>
17	consth	W	07F		<i>7Bitmask</i>
18	consth	W	0FFF		<i>12Bitmask</i>
19	consth	W	0FFFFF		<i>20Bitmask</i>
20	constd	W	7		<i>shiftToRd</i>
21	constd	W	15		<i>shiftToRs1</i>
22	constd	W	20		<i>shiftToRs2</i>
23	constd	W	12		<i>shiftToFunct3</i>
24	constd	W	25		<i>shiftToFunct7</i>
25	constd	W	5		<i>shiftBy5</i>
26	constd	W	11		<i>shiftBy11</i>
27	constd	W	3		<i>load</i>
28	constd	W	19		<i>opImm</i>
29	constd	W	23		<i>auipc</i>
30	constd	W	27		<i>opImm32</i>
31	constd	W	35		<i>store</i>
32	constd	W	51		<i>op</i>
33	constd	W	55		<i>lui</i>
34	constd	W	59		<i>op32</i>
35	constd	W	99		<i>branch</i>
36	constd	W	103		<i>jalr</i>
37	constd	W	111		<i>jal</i>

(TODO: Maybe neusortieren, andere constanten aufnehmen. Explain)

Figure 6: Sorts and non-progressive Constants for encoding RISC-V in BTOR2

```

for  $i$  from 0 to 31 do
  add to model:
   $n$   constd  W   $i$        $iConst$ 
end

```

Algorithm 1: progressive constants for encoding RISC-V in BTOR2

$(n + 0)$	state	D	$x0$	$(n + 17)$	state	D	$x17$
$(n + 1)$	state	D	$x1$	$(n + 18)$	state	D	$x18$
$(n + 2)$	state	D	$x2$	$(n + 19)$	state	D	$x19$
$(n + 3)$	state	D	$x3$	$(n + 20)$	state	D	$x20$
$(n + 4)$	state	D	$x4$	$(n + 21)$	state	D	$x21$
$(n + 5)$	state	D	$x5$	$(n + 22)$	state	D	$x22$
$(n + 6)$	state	D	$x6$	$(n + 23)$	state	D	$x23$
$(n + 7)$	state	D	$x7$	$(n + 24)$	state	D	$x24$
$(n + 8)$	state	D	$x8$	$(n + 25)$	state	D	$x25$
$(n + 9)$	state	D	$x9$	$(n + 26)$	state	D	$x26$
$(n + 10)$	state	D	$x10$	$(n + 27)$	state	D	$x27$
$(n + 11)$	state	D	$x11$	$(n + 28)$	state	D	$x28$
$(n + 12)$	state	D	$x12$	$(n + 29)$	state	D	$x29$
$(n + 13)$	state	D	$x13$	$(n + 30)$	state	D	$x30$
$(n + 14)$	state	D	$x14$	$(n + 31)$	state	D	$x31$
$(n + 15)$	state	D	$x15$	$(n + 32)$	state	AS	pc
$(n + 16)$	state	D	$x16$	$(n + 33)$	state	Mem	$memory$

Figure 7: State representation for encoding

4.2.2 State Representation

The next logical step is defining a representation of a RISC-V state. This is straightforward, as shown in Figure 7. I also introduced a flag for each register in my code to track whether the register was written to, enabling the transformation of a witness to a state file containing only the relevant registers. As these flags do not affect the operation of the BTOR2 model and are only included for an aesthetic choice, they are not included in my description and will not be discussed further.

4.2.3 Initialization

To initialize a state in BTOR2 from a RISC-V state file, the values in the registers must be loaded as constants, and for each memory address mentioned in the state file, the value and address must be loaded as constants. Due to the inability to represent a full 64-bit address space, I must manage the reduction of the address space from the state file to the BTOR2 model. I chose to initialize only the addresses up to the BTOR2 model's address space maximum and omit all others from the state file, as this provides the most predictable behavior. All addresses not mentioned in the state file are zero-initialized. Finally, these constants are used to initialize the state. For the registers, this is straightforward; for memory, all memory addresses are first written into a placeholder array, which is then used to initialize the actual memory. Due to BTOR2 constraints, these constants must be defined **before** the states, but initialization with the values must occur after the states. Thus, this initialization process **wraps around** the state representation. The generation of constants is shown in Algorithm 2, while the actual initialization is shown in Algorithm 3.

4.2.4 Fetching the Current Instruction

To fetch the current instruction, I read the four bytes of the instruction and concatenate them, as shown in Figure 8.

4.2.5 Deconstruction of the Instruction

With the instruction available, it can be deconstructed to extract the *opcode*, *rd*, *rs1*, *rs2*, *funct3*, *funct7*, and *imm*. For everything except *imm*, this can be accomplished by shifting and masking, as shown in Figure 9.

The immediate, however, must first be constructed from its subfields, which are referenced in Figure 1. In the BTOR2 model, this is shown in Figure 10. **(TODO:**

```

truePc ← value of pc in state file
maxPc ← number of addresses in BTOR2 model
pcValue ← truePc modulo maxPc
add to model:


---


n  constd  AS  pcValue  pcConst


---


for every register  $x_i$  do
  if register is initialised in state file then
    registerValue ← value of  $x_i$ 
    if registerValue ≠ 0 then
      add to model:
      

---


      n  constd  D  registerValue   $x_i$ Const
      

---


    end
  end
end
end
add to model:


---


(n + 0)  state  Mem  memPH
(n + 1)  init   Mem  memPH (n + 0)


---


lastPH ← memPH
allInitialCells ← all initialised memory cells in the state file
cutInitialCells ← remove all cells with address over maxPc
for every cell  $c$  in cutInitialCells do
  address ← address of  $c$ 
  value ← value of  $c$ 
  add to model:
  

---


  (n + 0)  constd  AS  address
  (n + 1)  constd  B   value
  (n + 2)  write   Mem  lastPH (n + 0) (n + 1) PHAfterC
  

---


  lastPH ← PHAfterC
end
keep lastPH for initialisation

```

Algorithm 2: Generating initialisation constants from state file in BTOR2

```

add to model:


---


n  init  AS  pc  pcConst


---


for every register  $x_i$  do
  if  $x_i$ Const was defined then
    add to model:
    

---


    n  init  D   $x_i$   $x_i$ Const
    

---


  end
end
end
add to model:


---


n  init  Mem  memory lastPh


---



```

Algorithm 3: Initialising states in the BTOR2 model

(n + 0)	read	B	memory	pc	instrB1
(n + 1)	add	AS	addressInc	pc	pc+1
(n + 2)	read	B	memory	pc+1	instrB2
(n + 3)	add	AS	addressInc	pc+1	pc+2
(n + 4)	read	B	memory	pc+2	instrB3
(n + 5)	add	AS	addressInc	pc+2	pc+3
(n + 6)	read	B	memory	pc+3	instrB4
(n + 7)	concat	H	instrB2	instrB1	instrH1
(n + 8)	concat	H	instrB4	instrB3	instrH2
(n + 9)	concat	W	instrH2	instrH1	instr

Figure 8: Fetching the current instruction from memory

(n + 0)	and	W	<i>instr</i>	<i>7Bitmask</i>	<i>opcode</i>
(n + 1)	srl	W	<i>instr</i>	<i>shiftToRd</i>	<i>rdPre</i>
(n + 2)	and	W	<i>rdPre</i>	<i>5Bitmask</i>	<i>rd</i>
(n + 3)	srl	W	<i>instr</i>	<i>shiftToRs1</i>	<i>rs1Pre</i>
(n + 4)	and	W	<i>rs1Pre</i>	<i>5Bitmask</i>	<i>rs1</i>
(n + 5)	srl	W	<i>instr</i>	<i>shiftToRs2</i>	<i>rs2Pre</i>
(n + 6)	and	W	<i>rs2Pre</i>	<i>5Bitmask</i>	<i>rs2</i>
(n + 7)	srl	W	<i>instr</i>	<i>shiftToFunct3</i>	<i>funct3Pre</i>
(n + 8)	and	W	<i>funct3Pre</i>	<i>shiftRd</i>	<i>funct3</i>
(n + 9)	srl	W	<i>instr</i>	<i>shiftToFunct7</i>	<i>funct7</i>

Figure 9: Extraction of values from the instruction without *imm*

Reference to same method in riscvsim) Three points are noteworthy:

First, some immediate subfields overlap exactly. This is utilized in lines (n + 1) with the overlap of *imm*[11 : 5] for I- and S-type, and (n + 21) with J- and B-types *imm*[10 : 5] overlap. Second, as described in Section 2.2, the immediate is always sign-extended. This is achieved using arithmetic right shifts, which perform sign extension and correctly position the highest immediate bit. Third, at line (n + 8), sign extension requires a right shift by 19. As this matches the opcode for arithmetic instructions with immediate, I reused this constant.

Now, *iTypeImm*, *sTypeImm*, *bTypeImm*, *uTypeImm*, and *jTypeImm* are available. However, it is preferable to have a single node *imm* referencing the immediate value regardless of instruction. This is accomplished in Figure 11, where booleans are defined to check all opcodes that are neither R-type nor I-type. Then, if-then-else nodes are chained to select instructions of J-type, U-type, B-type, or S-type. If the instruction is none of these, I default to I-type, as R-type does not use an immediate value. Finally, *imm* is extended to the 64-bit width required by RV64I.

At this stage, the values of the designated *rs1* and *rs2* registers can also be extracted. This is shown for *rs1* in Figure 4; the process is identical for *rs2*, with only the names changed. The starting equality comparisons can be omitted for *rs2*, as they are already defined for *rs1* and can be referenced.

(n + 0)	sra	W	<i>instr</i>	<i>shiftToRs2</i>	<i>iTypeImm</i>
(n + 1)	and	W	<i>iTypeImm</i>	<i>-5Bitmask</i>	<i>s[11:5]</i>
(n + 2)	add	W	<i>s[11:5]</i>	<i>rd</i>	<i>sTypeImm</i>
(n + 3)	and	W	<i>rd</i>	<i>-bitPicker</i>	<i>b[4:0]</i>
(n + 4)	and	W	<i>funct7</i>	<i>6Bitmask</i>	<i>b[10:5]Pre</i>
(n + 5)	sll	W	<i>b10:5Pre</i>	<i>shiftBy5</i>	<i>b[10:5]</i>
(n + 6)	and	W	<i>bitPicker</i>	<i>rd</i>	<i>b[11]Pre</i>
(n + 7)	sll	W	<i>b[11]Pre</i>	<i>shiftBy11</i>	<i>b[11]</i>
(n + 8)	sra	W	<i>instr</i>	<i>mathI</i>	<i>b[31:12]Pre</i>
(n + 9)	and	W	<i>b[31:12]Pre</i>	<i>12Bitmask</i>	<i>b[31:12]</i>
(n + 10)	add	W	<i>b[10:5]</i>	<i>b[4:0]</i>	<i>b[10:0]</i>
(n + 11)	add	W	<i>b[11]</i>	<i>b[10:0]</i>	<i>b[11:0]</i>
(n + 12)	add	W	<i>b[31:12]</i>	<i>b[11:0]</i>	<i>bTypeImm</i>
(n + 13)	and	W	<i>instr</i>	<i>-12Bitmask</i>	<i>uTypeImm</i>
(n + 14)	and	W	<i>rs2</i>	<i>-bitPicker</i>	<i>j[4:0]</i>
(n + 15)	and	W	<i>rs2</i>	<i>bitPicker</i>	<i>j[11]Pre</i>
(n + 16)	sll	W	<i>j[11]Pre</i>	<i>shiftBy11</i>	<i>j[11]</i>
(n + 17)	sll	W	<i>funct3</i>	<i>shiftToFunct3</i>	<i>j[14:12]</i>
(n + 18)	sll	W	<i>rs1</i>	<i>shiftToRs1</i>	<i>j[19:15]</i>
(n + 19)	sra	W	<i>instr</i>	<i>shiftBy11</i>	<i>j[31:20]Pre</i>
(n + 20)	and	W	<i>j[31:20]Pre</i>	<i>-20Bitmask</i>	<i>j[31:20]</i>
(n + 21)	add	W	<i>b[10:5]</i>	<i>j[4:0]</i>	<i>j[10:0]</i>
(n + 22)	add	W	<i>j[11]</i>	<i>j[10:0]</i>	<i>j[11:0]</i>
(n + 23)	add	W	<i>j[14:12]</i>	<i>j[11:0]</i>	<i>j[14:0]</i>
(n + 24)	add	W	<i>j[19:15]</i>	<i>j[14:0]</i>	<i>j[19:0]</i>
(n + 25)	add	W	<i>j[31:20]</i>	<i>j[19:0]</i>	<i>jTypeImm</i>

Figure 10: Extraction of all *imm* types from the instruction

(n + 0)	eq	Bool	opcode	store		isSType
(n + 1)	eq	Bool	opcode	branch		isBType
(n + 2)	eq	Bool	opcode	auipc		uType1
(n + 3)	eq	Bool	opcode	lui		uType2
(n + 4)	or	Bool	uType1	uType2		isUType
(n + 5)	eq	Bool	opcode	jal		isJType
(n + 6)	ite	W	isSType	sTypeImm	iTypeImm	checkS
(n + 7)	ite	W	isBType	bTypeImm	checkS	checkB
(n + 8)	ite	W	isUType	uTypeImm	checkB	checkU
(n + 9)	ite	W	isJType	jTypeImm	checkU	imm32
(n + 10)	sext	D	imm32	32		imm

Figure 11: Finding the correct immediate by opcode

```

for i from 1 to 31 do
  add to model:
    n  eq  Bool  rs1  iConst  isRs1Xi
end
add to model:
  n  ite  D  isRs1X1  x1  x0  checkX1
for i from 2 to 30 do
  add to model:
    n  ite  D  isRs1Xi  xi  checkX(i - 1)  checkXi
end
add to model:
  n  ite  D  isRs1X31  x31  checkX30  rs1val

```

Algorithm 4: Extracting the value of the register designated by *rs1*

<i>(isJALR already exists)</i>					
n	and	Bool	<i>isLoad</i>	<i>is5Funct3</i>	<i>isLHU</i>
(n + 0)	consth	W	<i>20</i>		<i>SUBWf7</i>
(n + 1)	eq	Bool	<i>funct7</i>	<i>SUBWf7</i>	<i>fitsF7SUBW</i>
(n + 2)	and	Bool	<i>is0Funct3</i>	<i>fitsF7SUBW</i>	<i>fitsF3SUBW</i>
(n + 3)	and	Bool	<i>isLoad</i>	<i>fitsF3SUBW</i>	<i>isSUBW</i>

(TODO: Use subfigs)

Figure 12: Instruction detection of JALR, LHU and SUBW as described in Algorithm 5

4.2.6 Instruction Detection

For the next-state logic, it is essential to determine the current command. Therefore, I defined a check *isInstruction* for each instruction. As this is repetitive, Algorithm 5 describes a generalized approach to obtain these booleans. An example for each instruction subgroup in Algorithm 5 is provided in Figure 12. The *funct7* checks from the *needsf7* subgroup can be reused if multiple instructions share the same *funct7*.

4.2.7 Next-State Logic

The next-state logic is the core of the model. Almost everything else supports this point. The goal is to create the changes each instruction would make and then apply only the changes specific to the instruction in the state. Each state node in the model must have an accompanying next node to function correctly. First, the changed values are computed.

Creating All Values of Instruction Execution

It is unnecessary to detail all instructions, as this simply follows the RV64I ISA. Instead, I provide examples for each group of instructions as divided in Table 1. Examples for AUIPC, JALR, BEQ, LHU, SD, ANDI, SLLIW, SLT, and SUBW are shown in Figure 13. These examples illustrate overlaps that can be utilized, such as addresses for load and store instructions or the 32-bit versions of word instructions. The SD

add to model:

(n + 0)	eq	Bool	opcode	load	isLoad
(n + 1)	eq	Bool	opcode	opImm	isOpImm
(n + 2)	eq	Bool	opcode	auipc	isAUIPC
(n + 3)	eq	Bool	opcode	opImm32	isOpImm32
(n + 4)	eq	Bool	opcode	store	isStore
(n + 5)	eq	Bool	opcode	op	isOp
(n + 6)	eq	Bool	opcode	lui	isLUI
(n + 7)	eq	Bool	opcode	op32	isOp32
(n + 8)	eq	Bool	opcode	branch	isBranch
(n + 9)	eq	Bool	opcode	jalr	isJALR
(n + 10)	eq	Bool	opcode	jal	isJAL

for i from 0 to 7 do

add to model:

n	eq	Bool	funct3	iConst	isiFunct3
---	----	------	--------	--------	-----------

end

onlyOp ← [LUI, AUIPC, JAL, JALR]

needsf7 ← [SRL, SRA, SRLI, SRAI, SRLW, SRAW, SRLWI, SRAWI, ADD, SUB, ADDW, SUBW]

rest ← [all other instructions]

for all instructions I in onlyOp do

isI is already defined

end

for all instructions I in rest do

opname ← opcode name of I

f3val ← expected funct3 of I as digit

add to model:

n	and	Bool	isopname	isf3valFunct3	isI
---	-----	------	----------	---------------	-----

end

for all instructions I in needsf7 do

opname ← opcode name of I

f3val ← expected funct3 of I as digit

f7hex ← expected funct7 of I as hexadecimal number

add to model:

(n + 0)	consth	W	f7hex	If7	
(n + 1)	eq	Bool	funct7	If7	fitsF7I
(n + 2)	and	Bool	isf3valFunct3	fitsF7I	fitsF3I
(n + 3)	and	Bool	isopname	fitsF3I	isI

end

Algorithm 5: Generalised approach to instruction detection

<i>n</i>	<i>add</i>	<i>D</i>	<i>imm</i>	<i>pc</i>	<i>rdAUIPC</i>
----------	------------	----------	------------	-----------	----------------

13.1: AUIPC

Figure 13: Instruction execution for chosen instructions

(<i>n</i> + 0)	<i>add</i>	<i>AS</i>	<i>pc</i>	<i>pcInc</i>	<i>nextPc</i>
(<i>n</i> + 1)	<i>add</i>	<i>D</i>	<i>imm</i>	<i>rs1val</i>	<i>pcJALR64pre</i>
(<i>n</i> + 2)	<i>and</i>	<i>D</i>	<i>-1Const</i>	<i>pcJALR64pre</i>	<i>pcJALR64</i>
(<i>n</i> + 3)	<i>slice</i>	<i>AS</i>	<i>pcJALR64</i>	<i>15</i>	<i>pcJALR</i>
(<i>n</i> + 4)	<i>uext</i>	<i>D</i>	<i>nextPc</i>	<i>48</i>	<i>rdJALR</i>

(TODO: pc overflow erwähnen)

13.2: JALR

Figure 13: Instruction execution for chosen instructions

example demonstrates that all other store instructions are interim results of preparing SD. Load instructions are similar, but each requires sign extension to 64 bits.

With this, each change can be assigned to its instruction.

The Next Memory

Defining the next memory array is straightforward. All store instructions are cascaded through if-then-else nodes, with the final 'else' set as the current memory array; if no 'if' matches, the array remains unchanged. This is shown in Figure 14.

(<i>n</i> + 0)	<i>add</i>	<i>AS</i>	<i>pc</i>	<i>pcInc</i>		<i>nextPc</i>
(<i>n</i> + 1)	<i>slice</i>	<i>AS</i>	<i>imm</i>	<i>15</i>	<i>0</i>	<i>ImmAS</i>
(<i>n</i> + 2)	<i>add</i>	<i>AS</i>	<i>pc</i>	<i>ImmAS</i>		<i>pcBranch</i>
(<i>n</i> + 3)	<i>eq</i>	<i>Bool</i>	<i>rs1val</i>	<i>rs2val</i>		<i>isBEQcond</i>
(<i>n</i> + 4)	<i>ite</i>	<i>AS</i>	<i>isBEQcond</i>	<i>pcBranch</i>	<i>nextPc</i>	<i>pcBEQ</i>

13.3: BEQ

Figure 13: Instruction execution for chosen instructions

(n + 0)	add	D	<i>rs1val</i>	<i>imm</i>		<i>1stAddrPre</i>
(n + 1)	slice	AS	<i>1stAddrPre</i>	<i>15</i>	<i>0</i>	<i>1stAddr</i>
(n + 2)	add	AS	<i>1stAddr</i>	<i>addressInc</i>		<i>2ndAddr</i>
(n + 3)	read	B	<i>memory</i>	<i>1stAddr</i>		<i>loadB1</i>
(n + 4)	read	B	<i>memory</i>	<i>2ndAddr</i>		<i>loadB2</i>
(n + 5)	concat	H	<i>loadB2</i>	<i>loadB1</i>		<i>loadB2B1</i>
(n + 6)	uext	D	<i>loadB2B1</i>	<i>48</i>	<i>0</i>	<i>rdLHU</i>

13.4: LHU

Figure 13: Instruction execution for chosen instructions

(n + 0)	add	D	<i>rs1val</i>	<i>imm</i>		<i>1stAddrPre</i>
(n + 1)	slice	AS	<i>1stAddrPre</i>	<i>15</i>	<i>0</i>	<i>1stAddr</i>
(n + 2)	add	AS	<i>1stAddr</i>	<i>addressInc</i>		<i>2ndAddr</i>
(n + 3)	add	AS	<i>2ndAddr</i>	<i>addressInc</i>		<i>3rdAddr</i>
(n + 4)	add	AS	<i>3rdAddr</i>	<i>addressInc</i>		<i>4thAddr</i>
(n + 5)	add	AS	<i>4thAddr</i>	<i>addressInc</i>		<i>5thAddr</i>
(n + 6)	add	AS	<i>5thAddr</i>	<i>addressInc</i>		<i>6thAddr</i>
(n + 7)	add	AS	<i>6thAddr</i>	<i>addressInc</i>		<i>7thAddr</i>
(n + 8)	add	AS	<i>7thAddr</i>	<i>addressInc</i>		<i>8thAddr</i>
(n + 9)	slice	B	<i>rs2val</i>	<i>7</i>	<i>0</i>	<i>storeB1</i>
(n + 10)	slice	B	<i>rs2val</i>	<i>15</i>	<i>8</i>	<i>storeB2</i>
(n + 11)	slice	B	<i>rs2val</i>	<i>23</i>	<i>16</i>	<i>storeB3</i>
(n + 12)	slice	B	<i>rs2val</i>	<i>31</i>	<i>24</i>	<i>storeB4</i>
(n + 13)	slice	B	<i>rs2val</i>	<i>39</i>	<i>32</i>	<i>storeB5</i>
(n + 14)	slice	B	<i>rs2val</i>	<i>47</i>	<i>40</i>	<i>storeB6</i>
(n + 15)	slice	B	<i>rs2val</i>	<i>55</i>	<i>48</i>	<i>storeB7</i>
(n + 16)	slice	B	<i>rs2val</i>	<i>63</i>	<i>56</i>	<i>storeB8</i>
(n + 17)	write	Mem	<i>memory</i>	<i>1stAddr</i>	<i>storeB1</i>	<i>memorySB</i>
(n + 18)	write	Mem	<i>memorySB</i>	<i>2ndAddr</i>	<i>storeB2</i>	<i>memorySH</i>
(n + 19)	write	Mem	<i>memorySH</i>	<i>3rdAddr</i>	<i>storeB3</i>	<i>memoryB3</i>
(n + 20)	write	Mem	<i>memoryB3</i>	<i>4thAddr</i>	<i>storeB4</i>	<i>memorySW</i>
(n + 21)	write	Mem	<i>memorySW</i>	<i>5thAddr</i>	<i>storeB5</i>	<i>memoryB5</i>
(n + 22)	write	Mem	<i>memoryB5</i>	<i>6thAddr</i>	<i>storeB6</i>	<i>memoryB6</i>
(n + 23)	write	Mem	<i>memoryB6</i>	<i>7thAddr</i>	<i>storeB7</i>	<i>memoryB7</i>
(n + 24)	write	Mem	<i>memoryB7</i>	<i>8thAddr</i>	<i>storeB8</i>	<i>memorySD</i>

13.5: SD

Figure 13: Instruction execution for chosen instructions

<i>n</i>	and	D	<i>rs1val</i>	<i>imm</i>	<i>rdANDI</i>
----------	-----	---	---------------	------------	---------------

13.6: ANDI

Figure 13: Instruction execution for chosen instructions

(<i>n</i> + 0)	and	W	<i>imm32</i>	<i>5Bitmask</i>	<i>shamtIW</i>
(<i>n</i> + 1)	slice	W	<i>rs1val</i>	<i>31</i>	<i>0</i> <i>rs1val32</i>
(<i>n</i> + 2)	sll	W	<i>rs1val32</i>	<i>shamtIW</i>	<i>rdSLLIWpre</i>
(<i>n</i> + 3)	sext	D	<i>rs1val32</i>	<i>32</i>	<i>rdSLLIW</i>

13.7: SLLIW

Figure 13: Instruction execution for chosen instructions

(<i>n</i> + 0)	slt	Bool	<i>rs1val</i>	<i>rs2val</i>	<i>rdSLTpre</i>
(<i>n</i> + 1)	uext	D	<i>rdSLTpre</i>	<i>63</i>	<i>rdSLT</i>

13.8: SLT

Figure 13: Instruction execution for chosen instructions

(<i>n</i> + 0)	slice	W	<i>rs1val</i>	<i>31</i>	<i>0</i> <i>rs1val32</i>
(<i>n</i> + 1)	slice	W	<i>rs2val</i>	<i>31</i>	<i>0</i> <i>rs2val32</i>
(<i>n</i> + 2)	sub	W	<i>rs1val32</i>	<i>rs2val32</i>	<i>rdSUBWpre</i>
(<i>n</i> + 3)	sext	D	<i>rdSUBWpre</i>	<i>32</i>	<i>rdSUBW</i>

13.9: SUBW

Figure 13: Instruction execution for chosen instructions

(<i>n</i> + 0)	ite	Mem	<i>isSB</i>	<i>memorySB</i>	<i>memory</i>	<i>newMem3</i>
(<i>n</i> + 1)	ite	Mem	<i>isSH</i>	<i>memorySH</i>	<i>newMem3</i>	<i>newMem2</i>
(<i>n</i> + 2)	ite	Mem	<i>isSW</i>	<i>memorySW</i>	<i>newMem2</i>	<i>newMem1</i>
(<i>n</i> + 3)	ite	Mem	<i>isSD</i>	<i>memorySD</i>	<i>newMem1</i>	<i>newMem</i>
(<i>n</i> + 4)	next	Mem	<i>memory</i>	<i>newMem</i>		

Figure 14: Next-State logic for the memory array

$(n + 0)$	ite	AS	<i>isBGEU</i>	<i>pcBGEU</i>	<i>nextPc</i>	<i>newPc7</i>
$(n + 1)$	ite	AS	<i>isBLTU</i>	<i>pcBLTU</i>	<i>newPc7</i>	<i>newPc6</i>
$(n + 2)$	ite	AS	<i>isBGE</i>	<i>pcBGE</i>	<i>newPc6</i>	<i>newPc5</i>
$(n + 3)$	ite	AS	<i>isBLT</i>	<i>pcBLT</i>	<i>newPc5</i>	<i>newPc4</i>
$(n + 4)$	ite	AS	<i>isBNE</i>	<i>pcBNE</i>	<i>newPc4</i>	<i>newPc3</i>
$(n + 5)$	ite	AS	<i>isBEQ</i>	<i>pcBEQ</i>	<i>newPc3</i>	<i>newPc2</i>
$(n + 6)$	ite	AS	<i>isJALR</i>	<i>pcJALR</i>	<i>newPc2</i>	<i>newPc1</i>
$(n + 7)$	ite	AS	<i>isJAL</i>	<i>pcJAL</i>	<i>newPc1</i>	<i>newPc</i>
$(n + 8)$	next	AS	<i>pc</i>	<i>newPc</i>		

Figure 15: Next-State logic for the pc register

The Next pc

For the next pc, the approach is similar, as shown in Figure 15. The only difference is that if no 'if' matches, pc must point to the next instruction to execute. The nextPc value was already computed for the JAL and JALR instructions and is reused here. The unconditional jumps also modify the value in rd, which is handled in the next section.

The Next rd

At last, the remaining registers must be updated. The procedure is defined in Figure 6. With the exception of x0, this is the same for all registers. The process is similar to defining the next memory or pc, but instead of a handful of instructions, all 39 relevant instructions must be considered, as only branch and store instructions do not modify rd. For brevity, the cascade for all relevant instructions is not shown in full in Algorithm 6, but only indicated.

4.2.8 Constraints

The final step is to define constraints to terminate the model checker. The primary constraint is reaching a set number of iterations, as shown in Figure 16.

add to model:							
n	next	D	<i>x0</i>	<i>x0</i>			
for i from 1 to 31 do							
add to model:							
(n + 0)	ite	D	<i>isLUI</i>	<i>rdLUI</i>	<i>xi</i>	<i>newXi-49</i>	
	:	ite	:	:	:	:	
(n + 47)	ite	D	<i>isSRAW</i>	<i>rdSRAW</i>	<i>newXi-2</i>	<i>newXi-1</i>	
(n + 48)	eq	Bool	<i>rd</i>	<i>iConst</i>		<i>isRdXi</i>	
(n + 49)	ite	D	<i>isRdXi</i>	<i>newXi-1</i>	<i>xi</i>	<i>newXi</i>	
(n + 50)	next	D	<i>xi</i>	<i>newXi</i>			
end							

Algorithm 6: Next-state logic for all x registers

(n + 0)	one	D			<i>counterInc</i>
(n + 1)	constd	D	<i>nIterations</i>	<i>maxIterations</i>	
(n + 2)	state	D			<i>counter</i>
(n + 3)	init	D	<i>counter</i>	<i>emptyReg</i>	
(n + 4)	add	D	<i>counter</i>	<i>counterInc</i>	<i>newCounter</i>
(n + 5)	next	D	<i>counter</i>	<i>newCounter</i>	
(n + 6)	eq	Bool	<i>counter</i>	<i>maxIterations</i>	<i>isMaxIter</i>
(n + 7)	bad		<i>isMaxIter</i>		

Figure 16: Constraining the model by iteration count

Additional constraints are defined to check for invalid instructions. The first checks if the opcode is valid for the model. The second constraint detects if the instruction cannot be identified even when the opcode is valid, as shown in Figure 17. The constraint in Figure 18 handles instruction-address-misaligned exceptions for jump instructions.

Other constraints can be defined, such as terminating on a specific pc value or when a register reaches a specified value.

(TODO: Maybe add examples on how to do this)

(TODO: Maybe internal references in figures should be numbers...)

(n + 0)	or	Bool	<i>isLoad</i>	<i>isOpImm</i>	<i>isOpcodeValid9</i>
(n + 1)	or	Bool	<i>isAUIPC</i>	<i>isOpcodeValid9</i>	<i>isOpcodeValid8</i>
(n + 2)	or	Bool	<i>isOpImm32</i>	<i>isOpcodeValid8</i>	<i>isOpcodeValid7</i>
(n + 3)	or	Bool	<i>isStore</i>	<i>isOpcodeValid7</i>	<i>isOpcodeValid6</i>
(n + 4)	or	Bool	<i>isOp</i>	<i>isOpcodeValid6</i>	<i>isOpcodeValid5</i>
(n + 5)	or	Bool	<i>isLUI</i>	<i>isOpcodeValid5</i>	<i>isOpcodeValid4</i>
(n + 6)	or	Bool	<i>isOp32</i>	<i>isOpcodeValid4</i>	<i>isOpcodeValid3</i>
(n + 7)	or	Bool	<i>isBranch</i>	<i>isOpcodeValid3</i>	<i>isOpcodeValid2</i>
(n + 8)	or	Bool	<i>isJALR</i>	<i>isOpcodeValid2</i>	<i>isOpcodeValid1</i>
(n + 9)	or	Bool	<i>isJAL</i>	<i>isOpcodeValid1</i>	<i>isOpcodeValid</i>
(n + 10)	bad		<i>-isOpcodeValid</i>		
(n + 11)	or	Bool	<i>isLUI</i>	<i>isAUIPC</i>	<i>isInstrValid47</i>
(n + 12)	or	Bool	<i>isJAL</i>	<i>isInstrValid47</i>	<i>isInstrValid46</i>
⋮	or	Bool	⋮	⋮	⋮
(n + 58)	or	Bool	<i>isSRAW</i>	<i>isInstrValid1</i>	<i>isInstrValid</i>
(n + 59)	and	Bool	<i>-isInstrValid</i>	<i>isOpcodeValid</i>	<i>unknownInstr</i>
(n + 60)	bad		<i>unknownInstr</i>		

Figure 17: Constraining the model on unknown instructions

(n + 0)	zero	AS			<i>pcZero</i>
(n + 1)	constd	AS	3		<i>pcBitmask</i>
(n + 2)	and	AS	<i>pcBitmask</i>	<i>pcJAL</i>	<i>lowbitsJAL</i>
(n + 3)	and	AS	<i>pcBitmask</i>	<i>pcJALR</i>	<i>lowbitsJALR</i>
(n + 4)	and	AS	<i>pcBitmask</i>	<i>pcBEQ</i>	<i>lowbitsBEQ</i>
⋮	and	AS	<i>pcBitmask</i>	⋮	⋮
(n + 9)	and	AS	<i>pcBitmask</i>	<i>pcBGEU</i>	<i>lowbitsBGEU</i>
(n + 10)	neq	Bool	<i>pcZero</i>	<i>lowbitsJAL</i>	<i>pcMsaJAL</i>
(n + 11)	neq	Bool	<i>pcZero</i>	<i>lowbitsJALR</i>	<i>pcMsaJALR</i>
(n + 12)	neq	Bool	<i>pcZero</i>	<i>lowbitsBEQ</i>	<i>pcMsaBEQ</i>
⋮	neq	Bool	<i>pcZero</i>	⋮	⋮
(n + 17)	neq	Bool	<i>pcZero</i>	<i>lowbitsBGEU</i>	<i>pcMsaBGEU</i>
(n + 18)	or	Bool	<i>pcMsaJAL</i>	<i>pcMsaJALR</i>	<i>pcMsa6</i>
(n + 19)	or	Bool	<i>pcMsaBEQ</i>	<i>pcBEQ</i>	<i>pcMsa5</i>
⋮	or	Bool	⋮	⋮	⋮
(n + 24)	or	Bool	<i>pcMsaBGEU</i>	<i>pcMsa1</i>	<i>pcMsa</i>
(n + 25)	bad		<i>pcMsa</i>		

Figure 18: Constraining the model on misaligned addresses

4.3 Testing for Correctness

To test my model, I compared its results to those of my RISC-V simulator (Section 2.3).

Given a state, both the simulation and the BTOR2 model are run with the iteration maximum set to 1. The resulting BTOR2 witness cannot be directly compared to the resulting state of the simulation. Therefore, I implemented a simple converter from witness to state [12, src/restate_witness.c]. These two states can then be compared. A shell script for this purpose is provided at [12, sh_utils/compare_iterations.sh].

To generate RISC-V states, I implemented a fuzzer [12, src/state_fuzzer.c] that generates randomized states with one valid instruction at the address of pc. The fuzzer first selects an instruction to test and fills all variable parts of the instruction, such as *rd* or *imm*. All registers relevant to the instruction are then assigned random 64-bit values. A pc value is generated to ensure the instruction fits within the limited address space of the BTOR2 model. If a jump instruction is chosen, possible address misalignment is corrected and address overflow is prevented. This simplifies later comparison of the resulting states, as correct execution of the instruction always results in the same state, despite differences between the simulation and the BTOR2 model.

With this setup, a series of tests can be conducted. For this, I implemented a shell script [12, sh_utils/test_btor2_model.sh]. As the number of tests increases, it becomes more challenging to track failed tests. To address this, I wrote a script to aggregate all failed tests into one file and add additional information such as instruction name or immediate value [12, sh_utils/diff_logger.sh].

I have executed approximately 5,000,000 tests on this model without a single failure, which leads me to conclude that my implementation is correct.

5 Benchmarks

With the model implemented, I was able to evaluate its performance. All benchmarks were executed on an Intel Core i5-6200U using the `btormc` model checker, which is distributed with the BTOR2 format [2]. Each test was run five times, and the resulting times were averaged. I also attempted to run the tests with the model checkers AVR [13] and Pono [14]. The challenges encountered with these tools are discussed in Section 5.2.

5.1 Tests

I devised two basic tests, each composed of four RISC-V instructions as illustrated in Figure 19. One test includes a memory operation, while the other does not, allowing for measurement of the impact of memory operations on the model’s performance. The programs for these tests are intentionally similar: both feature three instructions forming a loop and one instruction serving as a “workhorse”. The test names, `add` and `write_mem`, are derived from this key instruction. The program is embedded into a state, as exemplified in Figure 20, where the `add` program is configured for 256 loops. In this setup, `x1` acts as a loop limiter, `x2` as a loop counter, and `x3` as an accumulator. The instructions are placed in the initial bytes of memory. In contrast, the memory operation test uses `x2` as an address to store the first byte of a register, with `x3` serving as this register. Both tests were also implemented with increasing loop counts up to 2048 to provide a reference for processing more iterations.

bge x2 x1 0x10	jump out of program if $x1 = x2$
add x3 x3 x2	either (add counter onto x3)
sb x3 0x14(x2)	or (store the first byte of x3 at counter + 0x14)
addi x2 x2 0x1	increment counter in x2
jalr x0 x0 0x0	jump back to address 0

Figure 19: Base test cases for the benchmarks

REGISTERS:

PC:0

x1:100

x2:0

MEMORY:

0:001158E3 # BGE x2 x1 0x10

4:002181B3 # ADD x3 x3 x2

8:00110113 # ADDI x2 x2 1

c:00000067 # JALR x0 x0 0

Figure 20: Example state for benchmark add_0256

Additionally, I evaluated the impact of initialized memory on runtime . For this, I introduced tests with the prefix `fullmem`, where memory addresses 0x18 to 0xfff were filled with the bit pattern '0101'. The first four words of memory were filled with the test program. The fifth word could also be filled , but I opted to leave it zero-initialized to ensure the test terminates by jumping to this address and not finding a valid instruction. This guarantees termination. To further extend this evaluation, I added tests with double the initialized memory, filling up to address 0x1fff. **(TODO: benchmark them, add times)**

I also investigated the impact of address space size by generating models from `add_0256`, `add_1024`, `write_mem_0256`, and `write_mem_1024` with extended address space. These tests use the prefix “`extaddr_x`”, where `x` denotes the maximum address length in bits.

Furthermore, I implemented the base `add` tests in a manner similar to the BTOR2 model described by F. Schrögenderfer in his master thesis [3, Chapter 8]. These tests

use the prefix “nopc”, as the program counter is abstracted to activation flags for each instruction. The model generation is demonstrated using `nopc_add_0256` as an example. First, the necessary sorts and constants for RISC-V are defined:

```

1  sort    bitvec  1  bool
2  sort    bitvec  8  memcell
3  sort    bitvec 16  addressspace
4  sort    bitvec 64  register
5  sort    array   3 2 Memory
6  zero    1      false
7  one     1      true
8  zero    2      emptymem
9  zero    4      emptyreg
10 zero    3      pcinit
11 consth  3      4  pcinc
12 consth  3      10 instr1_pcmmod
13 consth  4      100 nloops

```

With this now the registers and memory can be defined:

```

99  state  3  pc      110 state  4  x10      121 state  4  x21
100 state  4  x0      111 state  4  x11      122 state  4  x22
101 state  4  x1      112 state  4  x12      123 state  4  x23
102 state  4  x2      113 state  4  x13      124 state  4  x24
103 state  4  x3      114 state  4  x14      125 state  4  x25
104 state  4  x4      115 state  4  x15      126 state  4  x26
105 state  4  x5      116 state  4  x16      127 state  4  x27
106 state  4  x6      117 state  4  x17      128 state  4  x28
107 state  4  x7      118 state  4  x18      129 state  4  x29
108 state  4  x8      119 state  4  x19      130 state  4  x30
109 state  4  x09     120 state  4  x20      131 state  4  x31

```

Note that node IDs 14-98 are skipped. As BTOR2 requires unique but not continuous node IDs, I assigned each register node ID to equal the register number plus 100 for clarity when writing models manually.

Next, the initial memory is defined:

144	state	5		169	write	5	166 167 168
145	init	5	14 8	170	consth	3	8
146	consth	3	0	171	consth	2	13
147	consth	2	e3	172	write	5	169 170 171
148	write	5	144 146 147	173	consth	3	9
149	consth	3	1	174	consth	2	01
150	consth	2	58	175	write	5	172 173 174
151	write	5	148 149 150	176	consth	3	a
152	consth	3	2	177	consth	2	11
153	consth	2	11	178	write	5	175 176 177
154	write	5	151 152 153	179	consth	3	b
155	consth	3	3	180	consth	2	00
156	consth	2	00	181	write	5	178 179 180
157	write	5	154 155 156	182	consth	3	c
158	consth	3	4	183	consth	2	67
159	consth	2	b3	184	write	5	181 182 183
160	write	5	157 158 159	185	consth	3	d
161	consth	3	5	186	consth	2	00
162	consth	2	81	187	write	5	184 185 186
163	write	5	16 161 162	188	consth	3	e
164	consth	3	6	189	consth	2	00
165	consth	2	21	190	write	5	187 188 189
166	write	5	16 164 165	191	consth	3	f
167	consth	3	7	192	consth	2	00
168	consth	2	00	193	write	5	190 191 192

Due to BTOR2 constraints, initial memory must be defined before the intended memory state node. Node 144 serves only for memory initialization, as previously described in Section 4.2.3.

The actual memory state and instruction flags are then defined. I renamed the flags from `stmt` (Schröngendorfer) to `instr` for clarity. An exit code is unnecessary, as differentiation between model terminations is not required; only the termination time is of interest.

199	state	5	<i>memory</i>
200	state	1	<i>instr0</i>
201	state	1	<i>instr1</i>
202	state	1	<i>instr2</i>
203	state	1	<i>instr3</i>
204	state	1	<i>endflag</i>

Next up would be the inputs and constraints for these. Both are not needed for the RISC-V model because neither are threads needed as I do not model parallel processing, nor is flushing as I use a naive memory model. And without inputs, the constraints are also nonexistent. So really, next up is initialization and transitions of the states. First is pc:

300	init	3	99	10		<i>pc is zero initialized</i>
301	add	3	99	11		<i>normal pc operation</i>
302	add	3	99	12		<i>instr1 jump</i>
303	eq	1	101	102		<i>instr1 branch condition</i>
304	ite	3	303	302	301	<i>IF condition THEN jump ELSE increment</i>
305	ite	3	200	304	99	<i>IF instr1 THEN check ELSE leave pc</i>
306	or	1	201	202		<i>instr2-3</i>
307	ite	3	306	301	305	<i>IF instr2-3 THEN increment ELSE instr1</i>
308	ite	3	203	10	307	<i>IF instr4 THEN jump0 ELSE try instr2-3</i>
309	next	3	99	308		

As the transition for pc are quite possibly the most complex I added some explanation.
x0 and x1 are skipped for now, so next registers x2 and x3:

314	init	4	102	9		319	init	4	103	9
315	one	4				320	add	4	102	103
316	add	4	102	315		321	ite	4	201	320
317	ite	4	202	316	102	322	next	4	103	321
318	next	4	102	317						

x2 increments by one each time the third instruction is run and x3 adds x2 every time the second instruction is run. All other registers and the memory do not change during execution, so I show them in one big block:

310	init	4	100	9	340	next	4	112	112	361	init	4	123	9
311	next	4	100	100	341	init	4	113	9	362	next	4	123	123
312	init	4	101	13	342	next	4	113	113	363	init	4	124	9
313	next	4	101	101	343	init	4	114	9	364	next	4	124	124
323	init	4	104	9	344	next	4	114	114	365	init	4	125	9
324	next	4	104	104	345	init	4	115	9	366	next	4	125	125
325	init	4	105	9	346	next	4	115	115	367	init	4	126	9
326	next	4	105	105	347	init	4	116	9	368	next	4	126	126
327	init	4	106	9	348	next	4	116	116	369	init	4	127	9
328	next	4	106	106	349	init	4	117	9	370	next	4	127	127
329	init	4	107	9	350	next	4	117	117	371	init	4	128	9
330	next	4	107	107	351	init	4	118	9	372	next	4	128	128
331	init	4	108	9	352	next	4	118	118	373	init	4	129	9
332	next	4	108	108	353	init	4	119	9	374	next	4	129	129
333	init	4	109	9	354	next	4	119	119	375	init	4	130	9
334	next	4	109	109	355	init	4	120	9	376	next	4	130	130
335	init	4	110	9	356	next	4	120	120	377	init	4	131	9
336	next	4	110	110	357	init	4	121	9	378	next	4	131	131
337	init	4	111	9	358	next	4	121	121	379	init	5	199	193
338	next	4	111	111	359	init	4	122	9	380	next	5	199	199
339	init	4	112	9	360	next	4	122	122					

Now only the instruction flags are left to handle:

381	init	1	200	7	<i>instr0 executes initally</i>
382	next	1	200	203	<i>instr0 only after instr3</i>
383	init	1	201	6	
384	and	1	200	-303	<i>instr0 and no branch</i>
385	next	1	201	384	<i>instr1 only after non-branching instr0</i>
386	init	1	202	6	
387	next	1	202	201	<i>instr2 only after instr2</i>
388	init	1	203	6	
389	next	1	203	202	<i>instr3 only after instr3</i>
390	init	1	204	6	
391	and	1	200	303	
392	next	1	204	391	<i>endflag if instr0 branches</i>
400	bad	204			<i>endflag terminates</i>

And with this, the model is can be run. The whole suite of add tests can be derived from this model by changing the constant value of node 13 to the appropriate loop count.

5.2 Results

To start I placed all iterations based benchmarks with btormc into Table 3 and the benchmarks of extended address space into Table 4. From Table 4 one can assume that the extension of address space does not impact the runtime of the model checker in a meaningful way. As seen in Table 5, with rising loop counts, memory operations need increasingly more time in comparison to their respective non-memory operation benchmarks. Also with rising loop counts the impact of large amounts of memory initialised is reduced, which was to be expected.

The `nopc` benchmark is significantly faster than my model. This was to be expected, as Schröndorfer models one RISC-V program specifically whereas I model the processor and feed it different programs by initialization. In this perspective, it is not surprising that a specialised model is faster than a generalised one. **(TODO: Add example for finding an instruction?)**

I also attempted to benchmark with the AVR and Pono model checkers, as they ranked first and second in the 2024 hardware model checking competition. Unfortunately, AVR did not terminate within 15 minutes when checking the `add_0256` benchmark.

loops	base		fullmem		nopc
	add	writemem	add	writemem	add
0256	2.635	2.877	9.759	13.344	0.136
0512	6.195	7.306	16.402	24.209	0.268
0768	10.802	13.283	24.093	36.388	0.414
1024	16.306	21.004	32.732	50.376	0.569
1280	23.032	30.200	42.410	65.746	0.728
1536	30.669	41.262	52.961	83.036	0.898
1792	39.463	53.940	64.598	101.475	1.075
2048	48.944	68.521	77.084	122.189	1.276

Table 3: Times of iterations based benchmarks

bits of address space	16	17	18	19	20
add_0256	2.635	2.632	2.626	2.626	2.624
add_1024	16.306	16.464	16.511	16.452	16.460
writemem_0256	2.877	2.88	2.890	2.889	2.890
writemem_1024	21.004	21.131	21.215	21.181	21.163

Table 4: Times of extended address space benchmarks

loops	$\frac{add}{writemem}$	$\frac{fullmem}{fullmem} \frac{add}{writemem}$	$\frac{add}{fullmem} \frac{add}{add}$	$\frac{writemem}{fullmem} \frac{writemem}{writemem}$	$\frac{add}{nopc} \frac{add}{add}$
0256	0.92	0.73	0.27	0.22	19.38
0512	0.85	0.68	0.38	0.3	23.12
0768	0.81	0.66	0.45	0.37	26.09
1024	0.78	0.65	0.5	0.42	28.66
1280	0.76	0.65	0.54	0.46	31.64
1536	0.74	0.64	0.58	0.5	34.15
1792	0.73	0.64	0.61	0.53	36.71
2048	0.71	0.63	0.63	0.56	38.36

Table 5: Relative runtime of benchmarks

Pono determined that this benchmark is satisfiable, but required nearly six minutes using the `-smt-solver cvc5` option. With `-smt-solver bzla`, Pono was slower but functional; other solvers reported inability to handle arrays. These results suggest that model checkers considered superior in general are not necessarily optimal for my model, whereas btormc, despite being unmaintained since August 2024, performed best. I suspect that newer model checkers are optimized for handling inputs, as most competition benchmarks involve at least one input beyond a clock signal, while my model operates solely with known constants and its execution time is dependent on rapid state iteration.

Another issue with AVR and Pono is their inability to generate a “complete” witness, as produced by the btormc option `-trace-gen-full`, which provides the values of all states in the final frame. Without this feature, it is not possible to reconstruct a state from their output after execution.

6 Conclusion

I developed tools to transition from a state to a model and from a witness of this model back to a state. Additionally, I implemented a fuzzer for the states to verify the correct functioning of the model, as well as a set of basic tests to benchmark its performance. Finally, I presented and discussed the results of these benchmarks, concluding that model checkers which appear superior in general are not necessarily better suited for this specific case.

ToDo Counters

To Dos: 18; 1, 2, ??, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18

Parts to extend: 1; 1

Draft parts: 0;

Bibliography

- [1] *The RISC-V Instruction Set Manual Volume I: Unprivileged ISA*, 2025, version 20250508. [Online]. Available: <https://lf-riscv.atlassian.net/wiki/spaces/HOME/pages/16154769/RISC-V+Technical+Specifications>
- [2] A. Niemetz, M. Preiner, C. Wolf, and A. Biere, “Btor2 , BtorMC and Boolector 3.0,” in *Computer Aided Verification*, H. Chockler and G. Weissenbacher, Eds. Cham: Springer International Publishing, 2018, pp. 587–595.
- [3] F. Schrögenderfer, “Bounded Model Checking of Lockless Programs,” Master’s thesis, Johannes Kepler University Linz, August 2021. [Online]. Available: <https://epub.jku.at/obvulihs/download/pdf/6579523>
- [4] A. Waterman, Y. Lee, D. A. Patterson, and K. Asanović, “The risc-v instruction set manual, volume i: Base user-level isa,” UC Berkeley, Tech. Rep. UCB/EECS-2011-62, May 2011. [Online]. Available: <http://www2.eecs.berkeley.edu/Pubs/TechRpts/2011/EECS-2011-62.html>
- [5] “History of RISC-V,” <https://riscv.org/about/>, accessed: 15.08.2025.
- [6] “RISC-V-Simulator,” <https://github.com/Kr1mo/Risc-V-Simulator>.
- [7] A. Biere, “Bounded model checking,” in *Handbook of Satisfiability*, 2nd ed., ser. Frontiers in Artificial Intelligence and Applications, A. Biere, M. Heule, H. van Maaren, and T. Walsh, Eds. IOS Press, 2021, vol. 336, pp. 739–764. [Online]. Available: <https://doi.org/10.3233/FAIA201002>

- [8] “HWMCC 2024,” <https://hwmcc.github.io/2024/>.
- [9] A. Biere, “The AIGER And-Inverter Graph (AIG) format version 20071012,” Institute for Formal Models and Verification, Johannes Kepler University, Altenbergerstr. 69, 4040 Linz, Austria, Tech. Rep. 07/1, 2007.
- [10] D. Beyer and J. Strejček, “Improvements in software verification and witness validation: SV-COMP 2025,” in *Proc. TACAS (3)*, ser. LNCS 15698. Springer, 2025, pp. 151–186.
- [11] D. Beyer, P.-C. Chien, and N.-Z. Lee, “Bridging hardware and software analysis with btor2c: A word-level-circuit-to-c translator,” in *Tools and Algorithms for the Construction and Analysis of Systems*, S. Sankaranarayanan and N. Sharygina, Eds. Cham: Springer Nature Switzerland, 2023, pp. 152–172.
- [12] “RISC-V_to_BTOR2,” https://github.com/Kr1mo/RISC-V_to_BTOR2.
- [13] A. Goel and K. Sakallah, “Avr: Abstractly verifying reachability,” in *Tools and Algorithms for the Construction and Analysis of Systems*, A. Biere and D. Parker, Eds. Cham: Springer International Publishing, 2020, pp. 413–422.
- [14] M. Mann, A. Irfan, F. Lonsing, Y. Yang, H. Zhang, K. Brown, A. Gupta, and C. Barrett, “Pono: A flexible and extensible smt-based model checker,” in *Computer Aided Verification: 33rd International Conference, CAV 2021, Virtual Event, July 20–23, 2021, Proceedings, Part II*. Berlin, Heidelberg: Springer-Verlag, 2021, p. 461–474. [Online]. Available: https://doi.org/10.1007/978-3-030-81688-9_22

(TODO: Add repo versions)