

Bachelor Thesis

Benchmark of RISC-V in BTOR2

Jan Krister Möller

Examiner: Dr. Mathias Fleury

University of Freiburg
Faculty of Engineering
Department of Computer Science
Chair of Computer Architecture

August 21, 2025

Writing Period

24. 06. 2025 – 24. 09. 2025

Examiner

Dr. Mathias Fleury

Declaration

I hereby declare that I am the sole author and composer of my thesis and that no other sources or learning aids, other than those listed, have been used. Furthermore, I declare that I have acknowledged the work of others by providing detailed references of said work.

I hereby also declare that my Thesis has not been prepared for another examination or assignment, either wholly or excerpts thereof.

Place, Date

Signature

Declaration on Usage of generative AI

I hereby declare that, with the approval of my examiner, I have employed the generative AI tool "GitHub Copilot" during the preparation of this thesis solely for spellchecking and enhancing the formality of my written expression. Furthermore, I expressly confirm that this tool was not used to generate data or as a source of factual information or content for this thesis.

Place, Date

Signature

Abstract

foo bar [1] [2] [3]

Contents

1	Motivation	1
2	RISC-V	3
2.1	Overview	3
2.2	The RV64I ISA	4
2.3	Simulation of RISC-V	6
2.3.1	Representing the State of a RISC-V Processor	6
2.3.2	Instruction detection	6
2.3.3	Instruction execution	6
2.3.4	Saving the State of a RISC-V Processor	6
3	BTOR2	9
3.1	Model Checking	9
3.2	The BTOR2 Language	9
3.3	The BTOR2 Witness	10
4	Transforming RISC-V to BTOR2	11
4.1	The Concept	11
4.2	Encoding	12
4.2.1	Constants	13
4.2.2	State Representation	13
4.2.3	Initialization	13
4.2.4	Fetching the current instruction	15

4.2.5	Deconstruction of the instruction	15
4.2.6	Command Detection	18
4.2.7	Next-State Logic	18
4.2.8	Constraints	18
4.3	Testing for Correctness	19
4.3.1	State Fuzzer	19
4.3.2	Automated Logging	19
4.4	Functional vs Relational Next-State Logic	19
5	Benchmarks	21
5.1	MultiAdd in Functional and Relational Next-State-Logic	21
5.2	Memory Operations	21
5.3	Results	21
	Bibliography	23

List of Figures

1	RV64I encoding formats	5
2	Construction of .state files	7
3	Constants for transforming RISC-V to BTOR2	14
4	State representation for transforming RISC-V to BTOR2	14
5	Fetching instruction	17
6	Extraction (without immediate)	18

List of Tables

1	RV64I Instruction Subset	6
---	------------------------------------	---

Liste der Algorithmen

1	Generating initialisation constants	16
2	Initialising states	17

1 Motivation

This is a template for an undergraduate or master's thesis. The first sections are concerned with the template itself. If this is your first thesis, consider reading.

2 RISC-V

As the first foundation for my benchmarks and, consequently, this thesis, I will discuss RISC-V and its operational principles.

2.1 Overview

RISC-V is an open-source instruction set architecture first published in May 2011 by A. Waterman et al. [4]. As indicated by its name, it is based on the RISC design philosophy. **(TODO: Explain RISC (compare wiki))** Since 2015, the development of RISC-V has been coordinated by the RISC-V International Association, a non-profit corporation based in Switzerland since 2020 [5]. Its objectives include providing an *open* ISA that is freely available to all, a *real* ISA suitable for native hardware implementation, and an ISA divided into a *small* base integer ISA usable independently, for example in educational contexts, with optional standard extensions to support general-purpose software development [1, Chapter 1].

Currently, RISC-V comprises four base ISAs: RV32I, RV64I, RV32E, and RV64E, which can be extended with one or more of the 47 ratified extension ISAs [1, Preface].

(EXTEND: Additional content may be required here) (TODO: Mention little endian?)

For the purposes of this work, I will focus on a subset of the RV64I ISA.

2.2 The RV64I ISA

RV64I is not overly complex, but its structure is essential for understanding the subsequent work presented in this thesis. Therefore, I will explain all elements relevant to my research.

RV64I features 32 64-bit registers, labeled $x0$ – $x31$, where $x0$ is hardwired to zero across all bits. Registers $x1$ – $x31$ are general-purpose and may be interpreted by various instructions as collections of booleans, two’s complement signed binary integers, or unsigned integers. Additionally, there is a register called *pc*, which serves as the program counter and holds the address of the current instruction [1, Chapters 4.1, 2.1].

In RV64I, memory addresses are 64 bits in size. As the memory model is defined to be single-byte addressable, the address space of RV64I encompasses 2^{64} bytes [1, Chapter 1.4].

Like nearly all standard ISAs of RISC-V, RV64I employs a standard instruction encoding length of 32 bits, or one *word*. Only the compressed extension C introduces instructions with a length of 16 bits [1, Chapter 1.5], which is not relevant for this discussion. All RV64I instructions are encoded in one of the six formats illustrated in Figure 1.

The design of these formats results in the following features:

- Due to RISC-V’s little-endian nature, the *opcode*, which encodes the general instruction, is always read first. Further specification of the instruction via *funct3* and *funct7* is consistently located at the same positions.
- If utilized by the instruction, the destination register *rd* and the source registers *rs1* and *rs2* are always found in the same locations, simplifying decoding.



Figure 1: RV64I encoding formats, used in [1, Chapter 2.3]

- The highest bit of the immediate value *imm* is always bit 31, making it straightforward to sign-extend the immediate value.

Note that each immediate subfield is labeled with its bit position within the immediate value. Immediate values are always sign-extended to 31 bits, and in the case of U-, B-, and J-type formats, the missing lower bits are filled with zeros.

The instructions relevant to my work are listed in Table 1

I have divided the instructions in Table 1 into nine groups based on their operations. LUI and AUIPC move a high immediate into *rd*; JA* instructions are unconditional jumps, and B* instructions are conditional jumps. L* instructions load sign-extended values from memory, either as Byte, Halfword, Word, or Doubleword lengths. Conversely, S* instructions write values of the specified length to memory. **(TODO: arithmetic)** Note that the suffix U denotes operations where values are processed as unsigned.

I left out FENCE, ECALL and EBREAK instructions as without I/O interaction or an environment like an OS or a debugger, these are not needed.

INSTR	TYPE	INSTR	TYPE	INSTR	TYPE	INSTR	TYPE
LUI	U	LW	I	XORI	I	SLT	I
AUIPC	U	LD	I	ORI	I	SLTU	I
JAL	J	LBU	I	ANDI	I	XOR	I
JALR	I	LHU	I	SLLI	I	OR	I
BEQ	B	LWU	I	SRLI	I	AND	I
BNE	B	SB	S	SRAI	I	SLL	I
BLT	B	SH	S	ADDIW	I	SRL	I
BGE	B	SW	S	SLLIW	I	SRA	I
BLTU	B	SD	S	SRLIW	I	ADDW	I
BGEU	B	ADDI	I	SRAIW	I	SLLW	I
LB	I	SLTI	I	ADD	I	SRLW	I
LH	I	SLTIU	I	SUB	I	SRAW	I

Table 1: Subset of RV64I instructions (TODO: Maybe rework, not happy yet)

2.3 Simulation of RISC-V

(TODO: This may be better placed in Chapter 4, but the state file is relevant here.)

2.3.1 Representing the State of a RISC-V Processor

2.3.2 Instruction detection

2.3.3 Instruction execution

2.3.4 Saving the State of a RISC-V Processor

To preserve the current state of a RISC-V processor, both the registers and memory must be stored. For this purpose, I have devised the format shown in Figure 2. The

```
1  REGISTERS:
2  PC: current pc in hex
3  x(0 - 31): value of register in hex
4
5  MEMORY:
6  (address in hex): byte, halfword, word or doubleword in hex
```

Figure 2: Construction of .state files

minimal file consists only of the two designators "REGISTERS:" and "MEMORY:", the current *pc*, and one empty line.

3 BTOR2

The second foundation of my benchmarks is BTOR2, a word-level model checking format published by A. Niemetz et al. [2].

3.1 Model Checking

(TODO: Write something about model checking...)

3.2 The BTOR2 Language

Generally in BTOR2, every line represents either a sort or a node, where normally the line number acts as an identifier. A sort behaves similar to a type as with it, either the length of a bitvector or the size of an array of bitvectors is defined. Nodes on the other hand represent a value of a defined sort and come as constants, operations or constraints. These values can later on be referenced by the node identifier, so the line number. The syntax of BTOR2 can be found at [2, figure 1] and corresponding operators in [2, table 1]

Key features of BTOR2 include its ability to operate sequentially, which makes the implementation of a RISC-V structure highly convenient. The main feature is the **state** operator, which defines a node that is sequentially updated. With an **init** node, this state can be assigned an initial value, and with **next**, the sequentially next

state can be defined. Finally, constraints can be used to specify endpoints for a model. These endpoints may indicate that something unintended has occurred or that the intended information has been found. In either case, the resulting model is provided as a witness.

3.3 The BTOR2 Witness

After receiving a witness, it must be interpreted. On the second line of a witness, the constraint that was triggered is specified. Subsequently, for each sequential iteration, the witness first presents—marked with $\#x$, where x is the iteration number—a representation of all states in the current iteration. Second, marked with $@x$, all inputs for the iteration are listed.

(TODO: Maybe a bit more, its a bit bare bones)

4 Transforming RISC-V to BTOR2

This chapter addresses the main problem of the thesis: transforming RISC-V code into the BTOR2 format for benchmarking purposes. My primary reference for this endeavor is F. Schrögender's master's thesis, "Bounded Model Checking in Lockless Programs"[3], in which he describes, among other topics, an encoding concept for a minimal machine in a multiprocessor context [3, Chapter 2] and two approaches to next-state logic: a functional [3, Chapter 6] and a relational [3, Chapter 7] approach. I will focus on the relational approach; a discussion of both approaches can be found in Section 4.4.

4.1 The Concept

To successfully execute a RISC-V instruction, three fundamental steps must occur in sequence:

- Fetch the current instruction from memory
- Identify the instruction
- Execute the instruction

Due to the fixed instruction length of RISC-V, as mentioned in Section 2.2, fetching the current instruction is straightforward. Ultimately, we want a node that retrieves a *word* from memory at the location specified by *pc*.

For basic identification, the *opcode* must be extracted and checked. Depending on the opcode, further distinctions between instructions require extracting and checking *funct3* and, if necessary, *funct7*. Ultimately, we want a node for each instruction, which holds a boolean value indicating whether this instruction was fetched.

To execute the instruction, we need to extract the values of the immediate *imm* and, if used, the registers *rs1* and *rs2*. All instructions only modify *rd*, *pc*, or memory. Therefore, the next-state logic can be generalized for these three cases.

Memory is only modified when a store instruction is identified. As all store instructions share the same type, computing the memory address is consistent across them. The final step is overwriting the memory at this address.

For the *pc*, except for jump commands, it always increments to point to the next instruction. The two unconditional jumps, **JAL** and **JALR**, must be handled separately. For branch instructions, after determining whether the relevant condition for the instruction holds, we can generalize, as all branch instructions execute the same operation from this point onward.

With *rd*, generalization across instructions is not feasible. However, we can generalize across all possible registers by adding a check in each register's update function to determine whether the register in question is *rd*.

4.2 Encoding

For better visualisation in the BTOR2 code I will mark all sort-ids in grey, all node-ids in red and all non-id numbers blue. As described in the BTOR2 syntax [2, Figure 1], each line can get an accompanying symbol. Sadly those cant be used as an alias to the line numbers, but for increased clarity, in the following figures I will use them as such aliases. With this I can also start each new figure with the relative line number

n and it makes it feasible to describe processes with algorithms. In the following, i will describe how I construct a BTOR2 model for a RISC-V state file.

4.2.1 Constants

First off, I added the sorts needed and some general purpose useful constants into the BTOR2 model as seen in Figure 3. Of note is the Representation of the memory as an array of addressable memory cells of each 1byte. Obviously, the set address space of 16bit is magnitudes away of the expected address space of 64bit, but representing a 64bit addressable memory with its resulting $2^{64}B \approx 18Exabyte$ is not implementable. Therefore, as I needed a feasible amount of memory space, I artificially chose a 16bit address space as a soft minimum. With $65kB$ and therefore programs with possibly > 10000 instructions I deemed this memory sufficient for most use cases. Despite this, the encoding is implemented in such a way that the address space can be altered with. (TODO: benchmark Auswirkungen von memory size)

4.2.2 State Representation

The next logical step is defining a representation of a RISC-V state. This is straightforward as shown in Figure 4. I also introduced a flag for each register in my code. They track if the register was written to and makes it possible to shorten a state file transformed from a witness to only the relevant registers. As they have no impact on the operation of the BTOR2 model, I will not mention them again.

4.2.3 Initialization

To initialize a state in BTOR2 from a RISC-V state file, the values in the registers must be loaded as constants, and for each memory address mentioned in the state file, the value and address has to be loaded as constants. Due to the inability to

1	sort	bitvec	1	<i>Bool</i>
2	sort	bitvec	16	<i>AS</i>
3	sort	bitvec	8	<i>B</i>
4	sort	bitvec	16	<i>H</i>
5	sort	bitvec	32	<i>W</i>
6	sort	bitvec	64	<i>D</i>
7	sort	array	2 3	<i>Mem</i>
8	zero	D		<i>empty_reg</i>
9	zero	B		<i>empty_cell</i>
10	constd	W	31	<i>register_bitmask</i>
11	constd	W	7	<i>shift_rd</i>
12	constd	W	15	<i>shift_rs1</i>
13	constd	W	20	<i>shift_rs2</i>
14	constd	W	12	<i>shift_funct3</i>
15	constd	W	25	<i>shift_funct7</i>
16	one	W		<i>bit_picker</i>
17	one	Bool		<i>true</i>
18	zero	Bool		<i>false</i>

(TODO: Maybe neusortieren, andere constanten aufnehmen)

Figure 3: Constants for encoding

(n + 0)	state	D	<i>x0</i>	(n + 17)	state	D	<i>x17</i>
(n + 1)	state	D	<i>x1</i>	(n + 18)	state	D	<i>x18</i>
(n + 2)	state	D	<i>x2</i>	(n + 19)	state	D	<i>x19</i>
(n + 3)	state	D	<i>x3</i>	(n + 20)	state	D	<i>x20</i>
(n + 4)	state	D	<i>x4</i>	(n + 21)	state	D	<i>x21</i>
(n + 5)	state	D	<i>x5</i>	(n + 22)	state	D	<i>x22</i>
(n + 6)	state	D	<i>x6</i>	(n + 23)	state	D	<i>x23</i>
(n + 7)	state	D	<i>x7</i>	(n + 24)	state	D	<i>x24</i>
(n + 8)	state	D	<i>x8</i>	(n + 25)	state	D	<i>x25</i>
(n + 9)	state	D	<i>x9</i>	(n + 26)	state	D	<i>x26</i>
(n + 10)	state	D	<i>x10</i>	(n + 27)	state	D	<i>x27</i>
(n + 11)	state	D	<i>x11</i>	(n + 28)	state	D	<i>x28</i>
(n + 12)	state	D	<i>x12</i>	(n + 29)	state	D	<i>x29</i>
(n + 13)	state	D	<i>x13</i>	(n + 30)	state	D	<i>x30</i>
(n + 14)	state	D	<i>x14</i>	(n + 31)	state	D	<i>x31</i>
(n + 15)	state	D	<i>x15</i>	(n + 32)	state	AS	<i>pc</i>
(n + 16)	state	D	<i>x16</i>	(n + 33)	state	Mem	<i>memory</i>

Figure 4: State representation for encoding

represent a full 64bit address space, the shrinking of the address space from state file to BTOR2 model must be handled. I decided to just initialise the addresses up to the BTOR2 model address space maximum and cut all others in the state file as I deem this the most predictable behaviour. Everything not mentioned in the state file will be zero-initialised. At last these constants must be used to initialise the state. For the registers this is straight forward, for the memory we must first write all memory addresses into a placeholder array which then we can use to initialise the real memory. Due to constraints in BTOR2, these constants have to be defined **before** the states, but initialisation with the values must happen after the states. This means that this initialisation process **wrappes around** the state representation. The generation of constants is shown in Algorithmus 1, whereas the actual initialization is shown in Algorithmus 2.

4.2.4 Fetching the current instruction

To fetch the current instruction, i read the 4 bytes of the instruction and concatenate them as seen in Figure 5

4.2.5 Deconstruction of the instruction

Now having the instruction, we can deconstruct it to extract the *opcode*, *rd*, *rs1*, *rs2*, *funct3*, *funct7* and *imm*. For everything apart of *imm*, this can be done by a shift and a masking. This is shown in Figure 6.

```

true_pc ← value of pc in state file
max_pc ← number of addresses in BTOR2 model
pc_value ← true_pc modulo max_pc
add to model:


---




---


(n + 0)  constd  AS  pc_value  pc_const


---


n += 1
for every register xi do
  if register is initialised in state file then
    register_value ← value of xi
    if register_value ≠ 0 then
      add to model:
      

---


      

---


      (n + 0)  constd  D  register_value  xi_const
      

---


      n += 1
    end
  end
end
end
add to model:


---




---


(n + 0)  state  Mem  mem_ph
(n + 1)  init   Mem  mem_ph (n + 0)


---




---


n += 2
last_ph ← mem_ph
all_initial_cells ← all initialised memory cells in the state file
cut_initial_cells ← remove all cells with address over max_pc
for every cell c in cut_initial_cells do
  address ← address of c
  value ← value of c
  add to model:
  

---


  

---


  (n + 0)  constd  AS  address
  (n + 1)  constd  B   value
  (n + 2)  write   Mem  last_ph (n + 0) (n + 1)  ph_after_c
  

---


  

---


  n += 3
  last_ph ← ph_after_c
end
keep last_ph for initialisation

```

Algorithmus 1 : Generating initialisation constants from state file in BTOR2


```

add to model:


---


(n + 0)  init  AS  pc  pc_const


---


n += 1

for every register  $x_i$  do
  if  $x_i\_const$  was defined then
    add to model:
    

---


    (n + 0)  init  D   $x_i$    $x_i\_const$ 
    

---


    n += 1
  end
end

add to model:


---


(n + 0)  init  Mem  memory  last_ph


---


n += 1

```

Algorithmus 2 : Initialising states in the BTOR2 model

(n + 0)	one	AS		<i>address_inc</i>
(n + 1)	read	B	<i>memory pc</i>	<i>instr_b1</i>
(n + 2)	add	AS	<i>address_inc pc</i>	<i>pc+1</i>
(n + 3)	read	B	<i>memory pc+1</i>	<i>instr_b2</i>
(n + 4)	add	AS	<i>address_inc pc+1</i>	<i>pc+2</i>
(n + 5)	read	B	<i>memory pc+2</i>	<i>instr_b3</i>
(n + 6)	add	AS	<i>address_inc pc+2</i>	<i>pc+3</i>
(n + 7)	read	B	<i>memory pc+3</i>	<i>instr_b4</i>
(n + 8)	concat	H	<i>instr_b2 instr_b1</i>	<i>instr_h1</i>
(n + 9)	concat	H	<i>instr_b4 instr_b3</i>	<i>instr_h2</i>
(n + 10)	concat	W	<i>instr_h2 instr_h1</i>	<i>instr</i>

(TODO: Maybe move constant to Figure 3)

Figure 5: Fetching the current instruction from memory

(n + 0)	consth	W	07f	opc_bitmask
(n + 0)	and	W	instr opc_bitmask	opcode
(n + 0)	srl	W	instr shift_rd	rd_pre
(n + 0)	and	W	rd_pre register_bitmask	rd
(n + 0)	srl	W	instr shift_rs1	rs1_pre
(n + 0)	and	W	rs1_pre register_bitmask	rs1
(n + 0)	srl	W	instr shift_rs2	rs2_pre
(n + 0)	and	W	rs2_pre register_bitmask	rs2
(n + 0)	srl	W	instr shift_funct3	funct3_pre
(n + 0)	and	W	funct3_pre shift_rd	funct3
(n + 0)	srl	W	instr shift_funct7	funct7

(TODO: Maybe move consth into constants Figure 3)

Figure 6: Extraction of *opcode*, *rd*, *rs1*, *rs2*, *funct3*, *funct7* and *imm* from the instruction

4.2.6 Command Detection

4.2.7 Next-State Logic

rd

pc

Memory

4.2.8 Constraints

(TODO: Iterations counter auch hier)

4.3 Testing for Correctness

4.3.1 State Fuzzer

4.3.2 Automated Logging

4.4 Functional vs Relational Next-State Logic

5 Benchmarks

5.1 MultiAdd in Functional and Relational Next-State-Logic

5.2 Memory Operations

5.3 Results

Bibliography

- [1] *The RISC-V Instruction Set Manual Volume I: Unprivileged ISA*, 2025, version 20250508. [Online]. Available: <https://lf-riscv.atlassian.net/wiki/spaces/HOME/pages/16154769/RISC-V+Technical+Specifications>
- [2] A. Niemetz, M. Preiner, C. Wolf, and A. Biere, “Btor2 , BtorMC and Boolector 3.0,” in *Computer Aided Verification*, H. Chockler and G. Weissenbacher, Eds. Cham: Springer International Publishing, 2018, pp. 587–595.
- [3] F. Schrögendorfer, “Bounded Model Checking of Lockless Programs,” Master’s thesis, Johannes Kepler University Linz, August 2021. [Online]. Available: <https://epub.jku.at/obvulihs/download/pdf/6579523>
- [4] A. Waterman, Y. Lee, D. A. Patterson, and K. Asanović, “The risc-v instruction set manual, volume i: Base user-level isa,” UC Berkeley, Tech. Rep. UCB/EECS-2011-62, May 2011. [Online]. Available: <http://www2.eecs.berkeley.edu/Pubs/TechRpts/2011/EECS-2011-62.html>
- [5] “History of RISC-V,” <https://riscv.org/about/>, accessed: 15.08.2025.

